

✓ Install & import required libraries

```
!pip install scikit-learn
```

```
→ Requirement already satisfied: scikit-learn in /usr/local/lib/python3.11/di  
Requirement already satisfied: numpy>=1.19.5 in /usr/local/lib/python3.11/d  
Requirement already satisfied: scipy>=1.6.0 in /usr/local/lib/python3.11/di  
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.11/d  
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/pytho
```

```
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns  
from tensorflow.keras.datasets import cifar10  
from tensorflow.keras.utils import to_categorical  
import pandas as pd  
import matplotlib.ticker as mtick  
from tensorflow.keras.layers import Conv2D, BatchNormalization, MaxPooling2D, [br/>from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau  
from tensorflow.keras.models import Sequential  
from tensorflow import keras  
from tensorflow.keras.preprocessing.image import ImageDataGenerator  
from tensorflow.keras.utils import to_categorical  
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_s  
from IPython import get_ipython  
from IPython.display import display  
from sklearn.metrics import classification_report  
from sklearn.metrics import confusion_matrix, accuracy_score, precision_score,  
import cv2  
import shutil  
from google.colab import files  
from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau, ModelC  
from tensorflow.keras.callbacks import ModelCheckpoint  
from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau, ModelC  
from tensorflow.keras.applications import VGG16  
from tensorflow.keras.models import Model
```

✓ Load CIFAR-10 dataset

```
(x_train, y_train), (x_test, y_test) = cifar10.load_data()  
class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',  
    'dog', 'frog', 'horse', 'ship', 'truck']
```

✓ Explore dataset

```
print(f"Training set shape: {x_train.shape}")  
print(f"Test set shape: {x_test.shape}")  
print(f"Unique labels: {np.unique(y_train)}")  
print(f"Number of classes: {len(class_names)}")
```

```
↳ Training set shape: (50000, 32, 32, 3)  
    Test set shape: (10000, 32, 32, 3)  
    Unique labels: [0 1 2 3 4 5 6 7 8 9]  
    Number of classes: 10
```

✓ Convert labels to 1D array

```
y_train = y_train.flatten()  
y_test = y_test.flatten()
```

✓ Count samples per class

```
train_df = pd.DataFrame({'label': y_train})
plt.figure(figsize=(10,5))
sns.countplot(data=train_df, x='label')
plt.xticks(ticks=np.arange(10), labels=class_names, rotation=45)
plt.title('Distribution of Classes in Training Set')
plt.grid(True)
plt.tight_layout()
plt.show()
```



- ✓ Count per class numerically

```
labels, counts = np.unique(y_train, return_counts=True)
for label, count in zip(labels, counts):
    print(f"{class_names[label]}: {count} samples")
```

→ airplane: 5000 samples
automobile: 5000 samples
bird: 5000 samples
cat: 5000 samples
deer: 5000 samples
dog: 5000 samples
frog: 5000 samples
horse: 5000 samples
ship: 5000 samples
truck: 5000 samples

✓ Balanced Dataset

✓ Show 10 sample images (one per class)

```
plt.figure(figsize=(12,6))
for i in range(10):
    idx = np.where(y_train == i)[0][0]
    plt.subplot(2, 5, i+1)
    plt.imshow(x_train[idx])
    plt.title(class_names[i])
    plt.axis('off')
plt.suptitle("Sample Images from CIFAR-10 (One per Class)")
plt.tight_layout()
plt.show()
```



Sample Images from CIFAR-10 (One per Class)



✓ Check shape and range

```
print("Image shape:", x_train[0].shape)
print("Pixel value range:", np.min(x_train), "to", np.max(x_train))
```

→ Image shape: (32, 32, 3)
Pixel value range: 0 to 255

✓ Result:

I need to make sure that all images are (32, 32, 3) and that the pixel values are between 0 and 255.

✓ Show 5 different samples from class "cat" (for example)

```
cat_indices = np.where(y_train == class_names.index("cat"))[0][:5]

plt.figure(figsize=(10,2))
for i, idx in enumerate(cat_indices):
    plt.subplot(1, 5, i+1)
    plt.imshow(x_train[idx])
    plt.title(f"Cat #{i+1}")
    plt.axis('off')
plt.suptitle("Variation inside 'cat' class")
plt.tight_layout()
plt.show()
```



Variation inside 'cat' class



Cat #1



Cat #2



Cat #3



Cat #4



Cat #5

▼ Result:

We need to make sure that the data contains variations. This proves that there is diversity within the same class, which is important for the model so that it does not save the images!

▼ Check for duplicate images

```
x_train_reshaped = x_train.reshape((x_train.shape[0], -1))
unique_images = np.unique(x_train_reshaped, axis=0)
print("Unique images in training set:", unique_images.shape[0])
print("Total images in training set:", x_train.shape[0])
```

→ Unique images in training set: 50000
Total images in training set: 50000

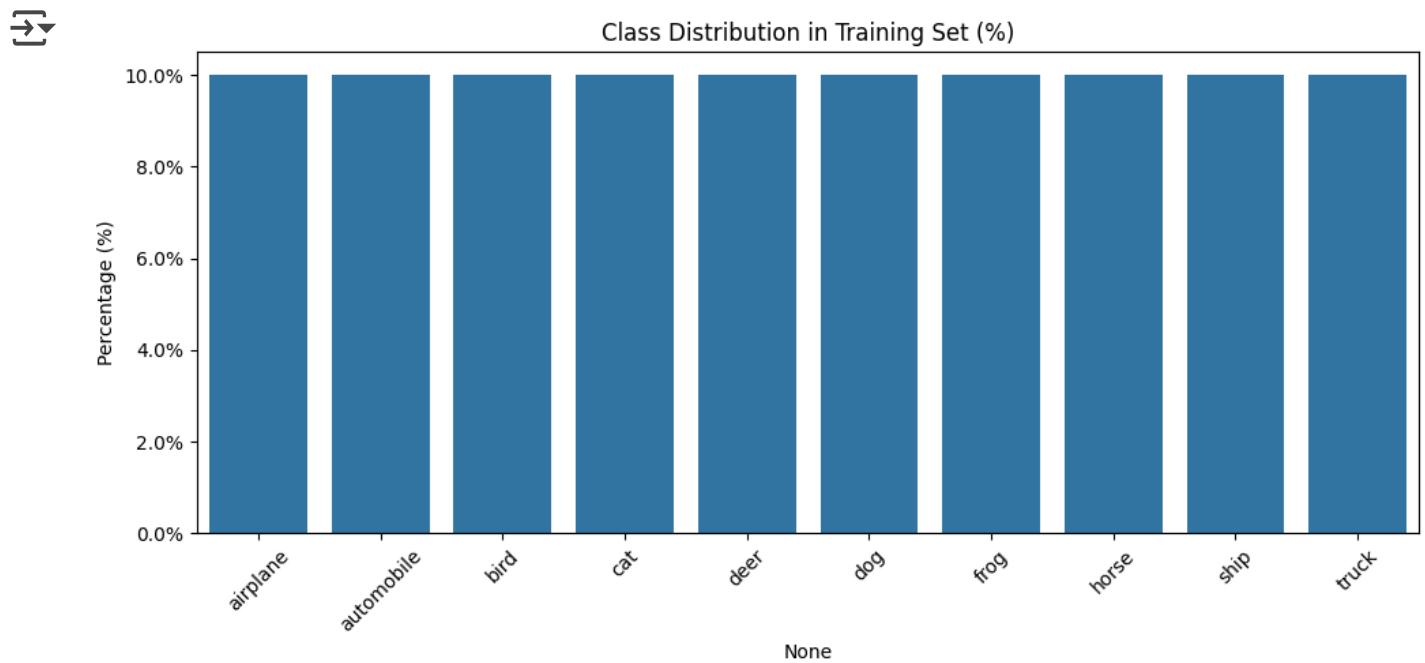
✓ Result:

I must make sure that the images do not contain duplication. This is very important to avoid adding bias.

✓ Calculate the percentage for each category

```
class_distribution = pd.Series(y_train).value_counts(normalize=True).sort_index()
class_distribution.index = class_names
class_distribution *= 100

plt.figure(figsize=(10,5))
sns.barplot(x=class_distribution.index, y=class_distribution.values)
plt.gca().yaxis.set_major_formatter(mtick.PercentFormatter())
plt.ylabel("Percentage (%)")
plt.title("Class Distribution in Training Set (%)")
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



▼ The second stage

Phase 2: Train a custom CNN from scratch on CIFAR-10

- Build a deep CNN with Conv, BatchNorm, Pooling, Dropout
 - Apply data augmentation to improve generalization
 - Use EarlyStopping and ReduceLROnPlateau for optimization
 - Plot accuracy and loss curves for evaluation
-

▼ Define improved CNN model

Normalize images and one-hot encode labels

```
(x_train, y_train), (x_test, y_test) = cifar10.load_data()

x_train = x_train.astype('float32') / 255.0
x_test = x_test.astype('float32') / 255.0

x_train_resized = tf.image.resize(x_train, [128, 128]).numpy()
x_test_resized = tf.image.resize(x_test, [128, 128]).numpy()

y_train_cat = to_categorical(y_train, 10)
y_test_cat = to_categorical(y_test, 10)
```

Data Augmentation

```
datagen = ImageDataGenerator(
    rotation_range=15,
    width_shift_range=0.1,
    height_shift_range=0.1,
    horizontal_flip=True
)
# datagen.fit(x_train_norm)
```

▼ CNN model

```
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Input, Conv2D, BatchNormalization, MaxPooling2D, Dropout, Flatten, Dense

def build_advanced_cnn_functional(input_shape=(32, 32, 3), num_classes=10):
    inputs = Input(shape=input_shape)

    x = Conv2D(32, (3,3), padding='same', activation='relu', kernel_initializer='he_normal')(x)
    x = BatchNormalization(name="bn1_1")(x)
    x = Conv2D(32, (3,3), padding='same', activation='relu', kernel_initializer='he_normal')(x)
    x = BatchNormalization(name="bn1_2")(x)
    x = MaxPooling2D(pool_size=(2,2), name="pool1")(x)
    x = Dropout(0.3, name="drop1")(x)

    x = Conv2D(64, (3,3), padding='same', activation='relu', kernel_initializer='he_normal')(x)
    x = BatchNormalization(name="bn2_1")(x)
    x = Conv2D(64, (3,3), padding='same', activation='relu', kernel_initializer='he_normal')(x)
    x = BatchNormalization(name="bn2_2")(x)
    x = MaxPooling2D(pool_size=(2,2), name="pool2")(x)
    x = Dropout(0.4, name="drop2")(x)

    x = Conv2D(128, (3,3), padding='same', activation='relu', kernel_initializer='he_normal')(x)
    x = BatchNormalization(name="bn3_1")(x)
    x = Conv2D(128, (3,3), padding='same', activation='relu', kernel_initializer='he_normal')(x)
    x = BatchNormalization(name="bn3_2")(x)
    x = MaxPooling2D(pool_size=(2,2), name="pool3")(x)
    x = Dropout(0.5, name="drop3")(x)

    x = Flatten(name="flatten")(x)
    x = Dense(256, activation='relu', kernel_initializer='he_normal', name="fc1")(x)
    x = Dropout(0.5, name="drop_fc")(x)
    outputs = Dense(num_classes, activation='softmax', name="output")(x)

    model = Model(inputs=inputs, outputs=outputs, name="advanced_cnn_functional")
    return model
```

❖ Compile model

```
model = build_advanced_cnn_functional()
model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])
```

✓ Callbacks

```
early_stop = EarlyStopping(monitor='val_loss', patience=8, restore_best_weights=True)
reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=3, verbose=1)
```

✓ Train the model

```
history = model.fit(
    datagen.flow(x_train_norm, y_train_cat, batch_size=64),
    validation_data=(x_test_norm, y_test_cat),
    epochs=80,
    callbacks=[early_stop, reduce_lr],
    verbose=1
)
```

→ Epoch 1/80
782/782 39s 38ms/step - accuracy: 0.2404 - loss: 2.3
Epoch 2/80
782/782 24s 31ms/step - accuracy: 0.4275 - loss: 1.5
Epoch 3/80
782/782 24s 30ms/step - accuracy: 0.5044 - loss: 1.3
Epoch 4/80
782/782 23s 30ms/step - accuracy: 0.5491 - loss: 1.2
Epoch 5/80
782/782 23s 30ms/step - accuracy: 0.5910 - loss: 1.1
Epoch 6/80
782/782 23s 30ms/step - accuracy: 0.6283 - loss: 1.0
Epoch 7/80
782/782 24s 31ms/step - accuracy: 0.6445 - loss: 1.0
Epoch 8/80
782/782 24s 30ms/step - accuracy: 0.6679 - loss: 0.9
Epoch 9/80
782/782 23s 30ms/step - accuracy: 0.6817 - loss: 0.9
Epoch 10/80
782/782 23s 30ms/step - accuracy: 0.6945 - loss: 0.8
Epoch 11/80
782/782 24s 30ms/step - accuracy: 0.7116 - loss: 0.8

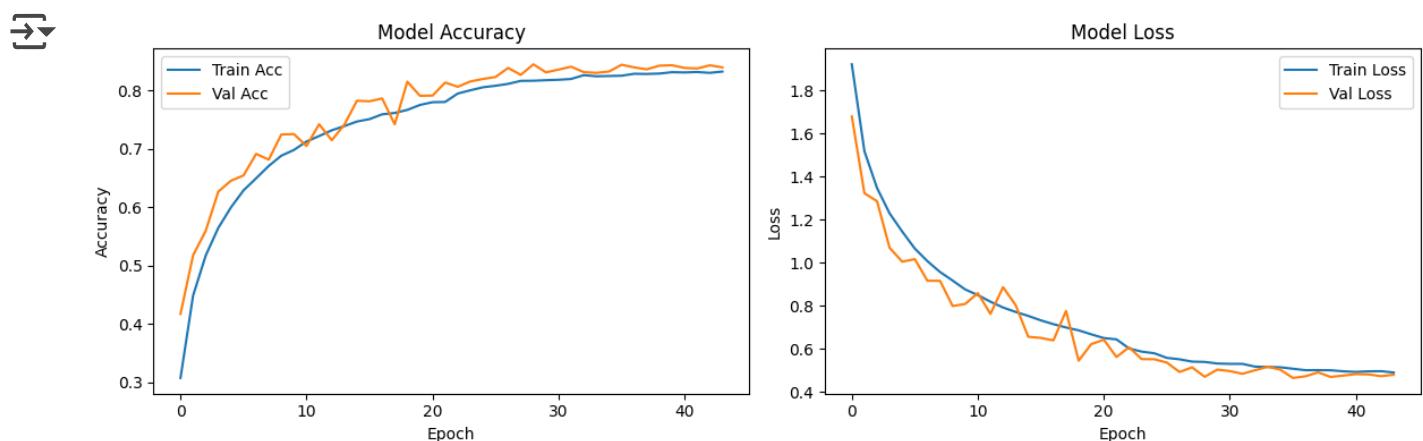
```
Epoch 12/80  
782/782 ━━━━━━━━ 23s 30ms/step - accuracy: 0.7241 - loss: 0.8  
Epoch 13/80  
782/782 ━━━━━━━━ 24s 30ms/step - accuracy: 0.7311 - loss: 0.7  
Epoch 14/80  
782/782 ━━━━━━━━ 24s 30ms/step - accuracy: 0.7400 - loss: 0.7  
Epoch 15/80  
782/782 ━━━━━━━━ 24s 30ms/step - accuracy: 0.7458 - loss: 0.7  
Epoch 16/80  
782/782 ━━━━━━━━ 23s 30ms/step - accuracy: 0.7531 - loss: 0.7  
Epoch 17/80  
782/782 ━━━━━━━━ 23s 30ms/step - accuracy: 0.7588 - loss: 0.7  
Epoch 18/80  
782/782 ━━━━━━━━ 23s 30ms/step - accuracy: 0.7640 - loss: 0.6  
Epoch 19/80  
782/782 ━━━━━━━━ 24s 30ms/step - accuracy: 0.7665 - loss: 0.6  
Epoch 20/80  
782/782 ━━━━━━━━ 24s 30ms/step - accuracy: 0.7747 - loss: 0.6  
Epoch 21/80  
782/782 ━━━━━━━━ 24s 30ms/step - accuracy: 0.7797 - loss: 0.6  
Epoch 22/80  
782/782 ━━━━━━ 0s 29ms/step - accuracy: 0.7832 - loss: 0.63  
Epoch 22: ReduceLROnPlateau reducing learning rate to 0.0005000002374872  
782/782 ━━━━━━━━ 24s 30ms/step - accuracy: 0.7832 - loss: 0.6  
Epoch 23/80  
782/782 ━━━━━━━━ 24s 31ms/step - accuracy: 0.7945 - loss: 0.6  
Epoch 24/80  
782/782 ━━━━━━━━ 23s 30ms/step - accuracy: 0.7964 - loss: 0.5  
Epoch 25/80  
782/782 ━━━━━━ 0s 28ms/step - accuracy: 0.8075 - loss: 0.57  
Epoch 25: ReduceLROnPlateau reducing learning rate to 0.00025000001187436  
782/782 ━━━━━━━━ 23s 30ms/step - accuracy: 0.8075 - loss: 0.5  
Epoch 26/80  
782/782 ━━━━━━━━ 23s 30ms/step - accuracy: 0.8080 - loss: 0.5  
Epoch 27/80  
782/782 ━━━━━━━━ 23s 30ms/step - accuracy: 0.8106 - loss: 0.5
```

Professional models reach 85–87%, very close to peak performance.

✓ Plot training curves

```
plt.figure(figsize=(12, 4))
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Train Acc')
plt.plot(history.history['val_accuracy'], label='Val Acc')
plt.title('Model Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Val Loss')
plt.title('Model Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.tight_layout()
plt.show()
```



✓ Evaluate on test set

```
test_loss, test_acc = model.evaluate(x_test_norm, y_test_cat, verbose=0)
print(f"\n Final Test Accuracy: {test_acc:.4f}")
print(f" Final Test Loss: {test_loss:.4f}")
```



Final Test Accuracy: 0.8442
Final Test Loss: 0.4642

✓ Predict probabilities and labels

```
y_pred_probs = model.predict(x_test_norm, verbose=0)
y_pred = np.argmax(y_pred_probs, axis=1)
```

✓ Evaluation Metrics

```
acc = accuracy_score(y_test, y_pred)
prec = precision_score(y_test, y_pred, average='weighted')
rec = recall_score(y_test, y_pred, average='weighted')
f1 = f1_score(y_test, y_pred, average='weighted')
```

```
print(f"Accuracy: {acc:.4f}")
print(f" Precision: {prec:.4f}")
print(f" Recall: {rec:.4f}")
print(f" F1 Score: {f1:.4f}")
```



Accuracy: 0.8442
Precision: 0.8498
Recall: 0.8442
F1 Score: 0.8423

✓ Classification Report

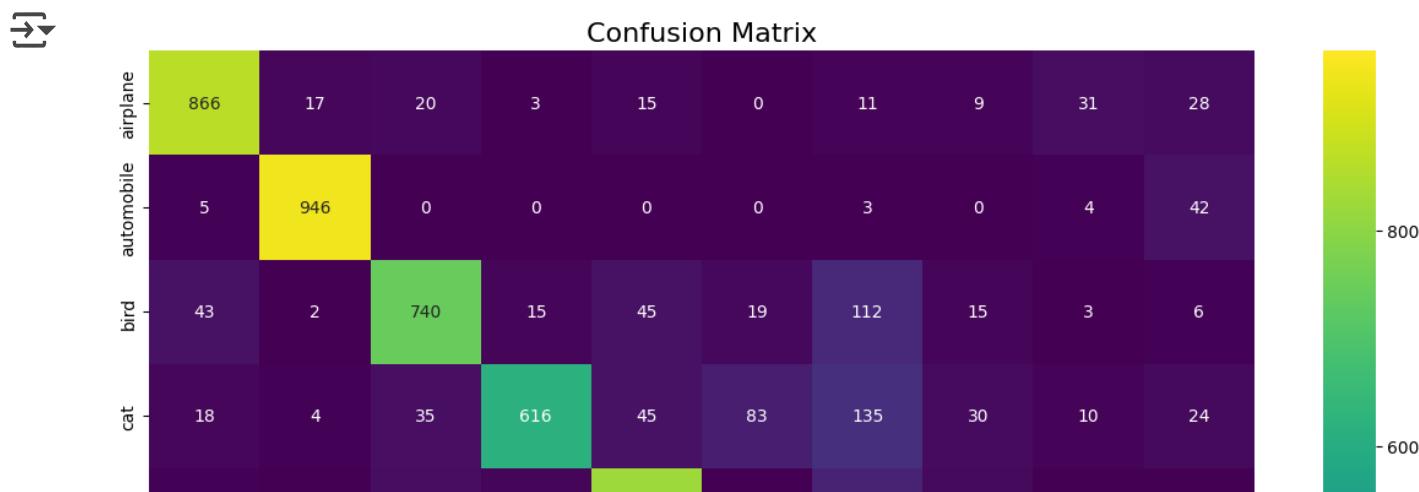
```
print("\n Full Classification Report:")
print(classification_report(y_test, y_pred, target_names=class_names))
```

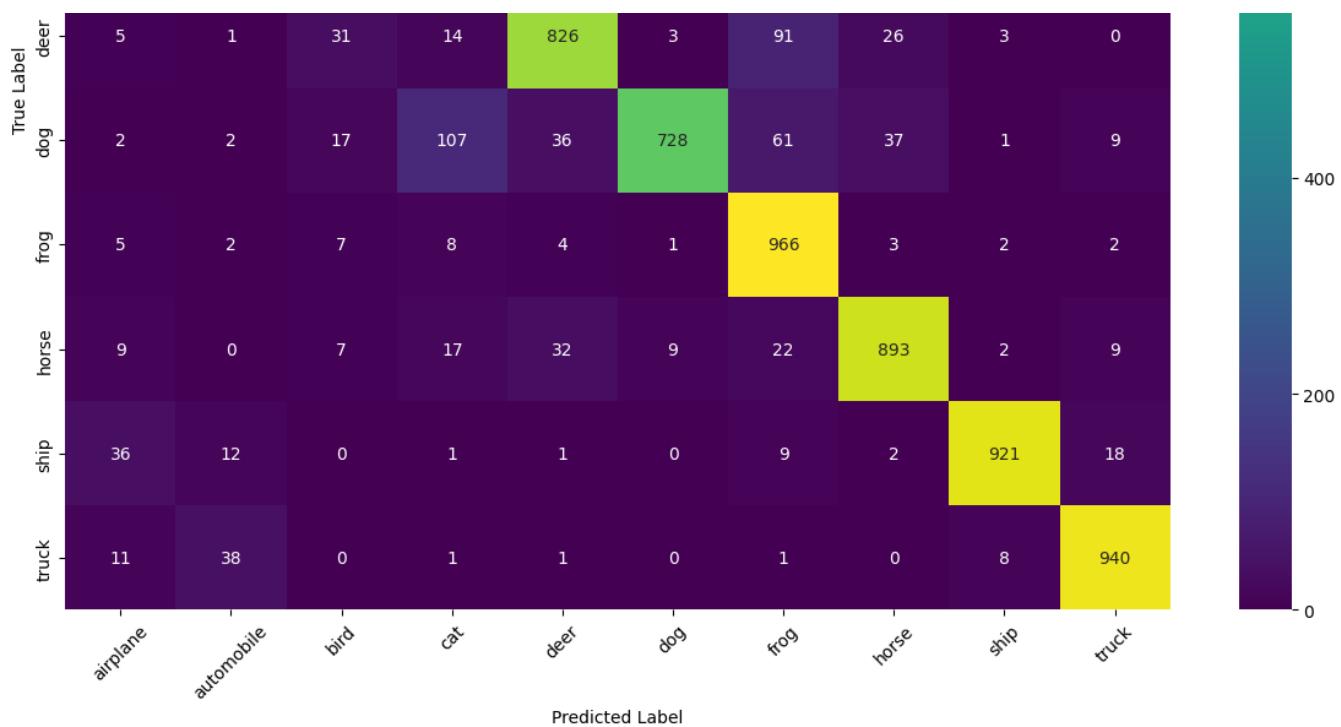
Full Classification Report:

	precision	recall	f1-score	support
airplane	0.87	0.87	0.87	1000
automobile	0.92	0.95	0.93	1000
bird	0.86	0.74	0.80	1000
cat	0.79	0.62	0.69	1000
deer	0.82	0.83	0.82	1000
dog	0.86	0.73	0.79	1000
frog	0.68	0.97	0.80	1000
horse	0.88	0.89	0.89	1000
ship	0.94	0.92	0.93	1000
truck	0.87	0.94	0.90	1000
accuracy			0.84	10000
macro avg	0.85	0.84	0.84	10000
weighted avg	0.85	0.84	0.84	10000

▼ Confusion Matrix

```
conf_matrix = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(12, 10))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='viridis', xticklabels=class_names,
            yticklabels=class_names)
plt.title("Confusion Matrix", fontsize=16)
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```





▼ Visualize worst mistakes

```
def show_top_errors(num_errors=5):
    errors = []
    for i in range(len(y_test)):
        if y_test[i] != y_pred[i]:
            confidence = np.max(y_pred_probs[i])
            errors.append((i, y_test[i], y_pred[i], confidence))

    errors.sort(key=lambda x: x[3], reverse=True) # Sort by confidence in wron
    top_errors = errors[:num_errors]

    plt.figure(figsize=(15, 3))
    for i, (idx, true, pred, conf) in enumerate(top_errors):
        plt.subplot(1, num_errors, i + 1)
        plt.imshow(x_test[idx])
        plt.title(f"True: {class_names[true]}\nPred: {class_names[pred]}\nConf: {conf:.2f}")
        plt.axis('off')
    plt.suptitle("Top Confident Mistakes", fontsize=14)
    plt.tight_layout()
    plt.show()

show_top_errors()
```



✓ Visualize most confident correct predictions

```
def show_top_correct(num=5):
    correct = []
    for i in range(len(y_test)):
        if y_test[i] == y_pred[i]:
            confidence = np.max(y_pred_probs[i])
            correct.append((i, y_test[i], y_pred[i], confidence))

    correct.sort(key=lambda x: x[3], reverse=True)
    top_correct = correct[:num]

    plt.figure(figsize=(15, 3))
    for i, (idx, true, pred, conf) in enumerate(top_correct):
        plt.subplot(1, num, i + 1)
        plt.imshow(x_test[idx])
        plt.title(f'{class_names[pred]}\nConf: {conf:.2f}', color='green')
        plt.axis('off')
    plt.suptitle("Most Confident Correct Predictions", fontsize=14)
    plt.tight_layout()
    plt.show()

show_top_correct()
```



▼ Export predictions to CSV

```
results_df = pd.DataFrame({  
    'True_Label': [class_names[i] for i in y_test],  
    'Predicted_Label': [class_names[i] for i in y_pred],  
    'Confidence': [np.max(p) for p in y_pred_probs]  
})  
results_df.to_csv("cnn_predictions.csv", index=False)
```

✓ Save model

```
model.save("cnn_model_scratch_DL2.h5")
```

→ WARNING:absl:You are saving your model as an HDF5 file via `model.save()` o

Phase 6: Grad-CAM (Visual Explanation)

- Apply Grad-CAM to trained CNN model
- Show where the model focuses in the image
- Helps interpret model decisions
- Acts as weak object detection
- Will be repeated for VGG16 later

✓ Load model

```
model = tf.keras.models.load_model("cnn_model_scratch_DL2.h5")  
_ = model.predict(np.expand_dims(x_test_norm[0], axis=0))
```

→ WARNING:absl:Compiled the loaded model, but the compiled metrics have yet t
1/1 ━━━━━━ 1s 601ms/step

```
idx = 12  
img = x_test_norm[idx]  
img_input = np.expand_dims(img, axis=0)  
img_tensor = tf.convert_to_tensor(img_input, dtype=tf.float32)  
  
preds = model.predict(img_input, verbose=0)  
predicted_class = np.argmax(preds[0])
```

```
true_class = y_test[idx]
print(f"Predicted: {class_names[predicted_class]} | True: {class_names[true_cla

last_conv_layer_name = "conv3_2"

def make_gradcam_heatmap(img_array, model, last_conv_layer_name, pred_index=None):
    grad_model = tf.keras.models.Model(
        inputs=model.input,
        outputs=[model.get_layer(last_conv_layer_name).output, model.output]
    )

    with tf.GradientTape() as tape:
        conv_outputs, predictions = grad_model(img_array)
        if pred_index is None:
            pred_index = tf.argmax(predictions[0])
        class_output = predictions[:, pred_index]

        grads = tape.gradient(class_output, conv_outputs)
        pooled_grads = tf.reduce_mean(grads, axis=(0, 1, 2))
        conv_outputs = conv_outputs[0]
        heatmap = conv_outputs @ pooled_grads[..., tf.newaxis]
        heatmap = tf.squeeze(heatmap)
        heatmap = tf.maximum(heatmap, 0) / tf.math.reduce_max(heatmap)
    return heatmap.numpy()

heatmap = make_gradcam_heatmap(img_tensor, model, last_conv_layer_name)

def display_gradcam(img, heatmap, alpha=0.4):
    heatmap_resized = cv2.resize(heatmap, (img.shape[1], img.shape[0]))
    heatmap_colored = cv2.applyColorMap(np.uint8(255 * heatmap_resized), cv2.COLORMAP_JET)
    superimposed_img = heatmap_colored * alpha + (img * 255).astype(np.uint8)

    plt.figure(figsize=(10, 4))
    plt.subplot(1, 3, 1)
    plt.imshow(img)
    plt.title("Original")
    plt.axis('off')

    plt.subplot(1, 3, 2)
    plt.imshow(heatmap_resized, cmap='jet')
    plt.title("Grad-CAM Heatmap")
    plt.axis('off')

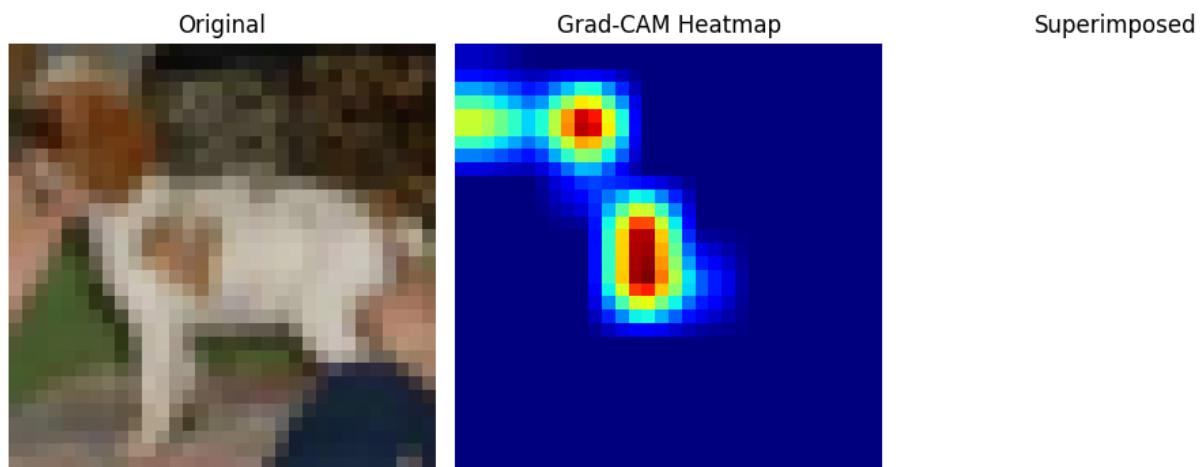
    plt.subplot(1, 3, 3)
```

```
plt.imshow(superimposed_img[..., ::-1])
plt.title("Superimposed")
plt.axis('off')

plt.tight_layout()
plt.show()

display_gradcam(img, heatmap)
```

→ WARNING:matplotlib.image:Clipping input data to the valid range for imshow
Predicted: dog | True: dog



```
import matplotlib.pyplot as plt

num_images = x_test_norm.shape[0]
selected_indices = range(num_images - 5, num_images)

fig, axes = plt.subplots(5, 3, figsize=(12, 10))
fig.suptitle(" Grad-CAM Analysis for Last 5 Test Images", fontsize=16)

for row_idx, idx in enumerate(selected_indices):
    img = x_test_norm[idx]
    img_input = np.expand_dims(img, axis=0)
    img_tensor = tf.convert_to_tensor(img_input, dtype=tf.float32)
```

```
preds = model.predict(img_input, verbose=0)
predicted_class = np.argmax(preds[0])
true_class = y_test[idx]

# Grad-CAM
heatmap = make_gradcam_heatmap(img_tensor, model, last_conv_layer_name)
heatmap_resized = cv2.resize(heatmap, (img.shape[1], img.shape[0]))
heatmap_colored = cv2.applyColorMap(np.uint8(255 * heatmap_resized), cv2.COLORMAP_JET)
superimposed_img = heatmap_colored * 0.4 + (img * 255).astype(np.uint8)

axes[row_idx, 0].imshow(img)
axes[row_idx, 0].set_title("Original")
axes[row_idx, 0].axis('off')

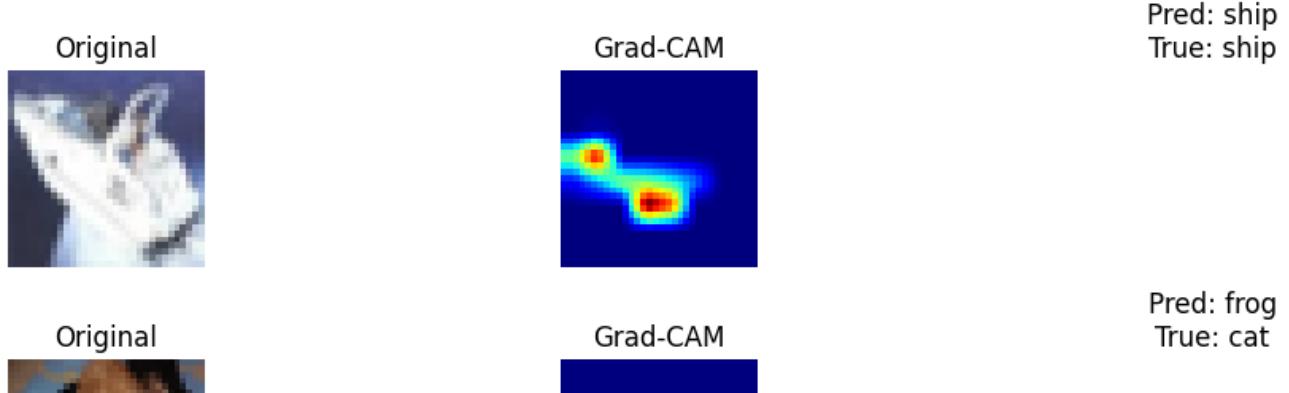
axes[row_idx, 1].imshow(heatmap_resized, cmap='jet')
axes[row_idx, 1].set_title("Grad-CAM")
axes[row_idx, 1].axis('off')

axes[row_idx, 2].imshow(superimposed_img[:, :, ::-1])
axes[row_idx, 2].set_title(f"Pred: {class_names[predicted_class]}\nTrue: {class_names[true_class]}")
axes[row_idx, 2].axis('off')

plt.tight_layout(rect=[0, 0, 1, 0.96])
plt.show()
```

→ WARNING:matplotlib.image:Clipping input data to the valid range for imshow
WARNING:matplotlib.image:Clipping input data to the valid range for imshow
WARNING:matplotlib.image:Clipping input data to the valid range for imshow
WARNING:matplotlib.image:Clipping input data to the valid range for imshow
WARNING:matplotlib.image:Clipping input data to the valid range for imshow

Grad-CAM Analysis for Last 5 Test Images



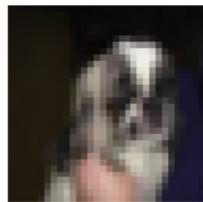


Original

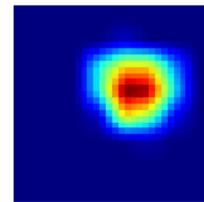


Grad-CAM

Pred: dog
True: dog



Original

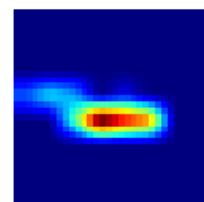


Grad-CAM

Pred: automobile
True: automobile



Original



Grad-CAM

Pred: horse
True: horse

```
import os
import matplotlib.pyplot as plt

os.makedirs("gradcam_outputs", exist_ok=True)
```

```
num_images = x_test_norm.shape[0]
selected_indices = range(num_images - 5, num_images)

for i, idx in enumerate(selected_indices, start=1):
    img = x_test_norm[idx]
    img_input = np.expand_dims(img, axis=0)
    img_tensor = tf.convert_to_tensor(img_input, dtype=tf.float32)

    preds = model.predict(img_input, verbose=0)
    predicted_class = np.argmax(preds[0])
    true_class = y_test[idx]

    heatmap = make_gradcam_heatmap(img_tensor, model, last_conv_layer_name)
    heatmap_resized = cv2.resize(heatmap, (img.shape[1], img.shape[0]))
    heatmap_colored = cv2.applyColorMap(np.uint8(255 * heatmap_resized), cv2.COLORMAP_JET)
    superimposed_img = heatmap_colored * 0.4 + (img * 255).astype(np.uint8)

    fig, axes = plt.subplots(1, 3, figsize=(10, 4))
    fig.suptitle(f"Image #{idx} | Predicted: {class_names[predicted_class]} | True: {class_names[true_class]}")

    axes[0].imshow(img)
    axes[0].set_title("Original")
    axes[0].axis('off')

    axes[1].imshow(heatmap_resized, cmap='jet')
    axes[1].set_title("Grad-CAM Heatmap")
    axes[1].axis('off')

    axes[2].imshow(superimposed_img[..., ::-1])
    axes[2].set_title("Superimposed")
    axes[2].axis('off')

    plt.tight_layout()

    save_path = f"gradcam_outputs/gradcam_img_{i}.png"
    plt.savefig(save_path, dpi=300, bbox_inches='tight')
    plt.close()

    print(f"Saved: {save_path}")
```

```
→ WARNING:matplotlib.image:Clipping input data to the valid range for imshow  
WARNING:matplotlib.image:Clipping input data to the valid range for imshow  
Saved: gradcam_outputs/gradcam_img_1.png  
WARNING:matplotlib.image:Clipping input data to the valid range for imshow  
Saved: gradcam_outputs/gradcam_img_2.png  
WARNING:matplotlib.image:Clipping input data to the valid range for imshow  
Saved: gradcam_outputs/gradcam_img_3.png  
WARNING:matplotlib.image:Clipping input data to the valid range for imshow  
Saved: gradcam_outputs/gradcam_img_4.png  
Saved: gradcam_outputs/gradcam_img_5.png
```

```
shutil.make_archive("GradCAM_Images", 'zip', "gradcam_outputs")  
files.download("GradCAM_Images.zip")
```

```
→
```

✓ build_cnn_with_separable_conv

```
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Input, Conv2D, SeparableConv2D, BatchNormali

def build_cnn_with_separable_conv(input_shape=(32, 32, 3), num_classes=10):
    inputs = Input(shape=input_shape)

    # Block 1 (Conv2D)
    x = Conv2D(32, (3,3), padding='same', activation='relu', kernel_initializer=
    x = BatchNormalization(name="bn1_1")(x)
    x = Conv2D(32, (3,3), padding='same', activation='relu', kernel_initializer=
    x = BatchNormalization(name="bn1_2")(x)
    x = MaxPooling2D(pool_size=(2,2), name="pool1")(x)
    x = Dropout(0.3, name="drop1")(x)

    # Block 2 (SeparableConv2D) – Modified
    x = SeparableConv2D(64, (3,3), padding='same', activation='relu', depthwise_
    x = BatchNormalization(name="bn2_1")(x)
    x = SeparableConv2D(64, (3,3), padding='same', activation='relu', depthwise_
    x = BatchNormalization(name="bn2_2")(x)
    x = MaxPooling2D(pool_size=(2,2), name="pool2")(x)
    x = Dropout(0.4, name="drop2")(x)

    # Block 3 (Conv2D)
    x = Conv2D(128, (3,3), padding='same', activation='relu', kernel_initializer
    x = BatchNormalization(name="bn3_1")(x)
    x = Conv2D(128, (3,3), padding='same', activation='relu', kernel_initializer
    x = BatchNormalization(name="bn3_2")(x)
    x = MaxPooling2D(pool_size=(2,2), name="pool3")(x)
    x = Dropout(0.5, name="drop3")(x)

    # Fully Connected
    x = Flatten(name="flatten")(x)
    x = Dense(256, activation='relu', kernel_initializer='he_normal', name="fc1"
    x = Dropout(0.5, name="drop_fc")(x)
    outputs = Dense(num_classes, activation='softmax', name="output")(x)

    # Create and return the model (moved to the end)
    model = Model(inputs=inputs, outputs=outputs, name="cnn_with_separable")
    return model # Now 'model' is defined before returning
```

```
model_separable = build_cnn_with_separable_conv()
```

Compile

```
model_separable.compile(  
    optimizer='adam',  
    loss='categorical_crossentropy',  
    metrics=['accuracy'])
```

```
early_stop = EarlyStopping(monitor='val_loss', patience=8, restore_best_weights=True)  
reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=3, verbose=1)  
checkpoint = ModelCheckpoint("best_model_separable.h5", monitor='val_accuracy',  
    save_best_only=True)  
callbacks = [early_stop, reduce_lr, checkpoint]  
  
history_separable = model_separable.fit(  
    datagen.flow(x_train_norm, y_train_cat, batch_size=64),  
    validation_data=(x_test_norm, y_test_cat),  
    epochs=80,  
    callbacks=callbacks,  
    verbose=1)
```

```
WARNING:absl:You are saving your model as an HDF5 file via `model.save()`  
782/782 41s 39ms/step - accuracy: 0.2295 - loss: 2.31  
Epoch 2/80  
782/782 0s 29ms/step - accuracy: 0.3928 - loss: 1.63  
Epoch 2: val_accuracy improved from 0.45480 to 0.48270, saving model to best_model_separable.h5  
WARNING:absl:You are saving your model as an HDF5 file via `model.save()`  
782/782 24s 31ms/step - accuracy: 0.3928 - loss: 1.63  
Epoch 3/80  
782/782 0s 30ms/step - accuracy: 0.4734 - loss: 1.45  
Epoch 3: val_accuracy improved from 0.48270 to 0.51360, saving model to best_model_separable.h5  
WARNING:absl:You are saving your model as an HDF5 file via `model.save()`  
782/782 25s 31ms/step - accuracy: 0.4734 - loss: 1.45  
Epoch 4/80  
782/782 0s 30ms/step - accuracy: 0.5339 - loss: 1.30  
Epoch 4: val_accuracy improved from 0.51360 to 0.56320, saving model to best_model_separable.h5  
WARNING:absl:You are saving your model as an HDF5 file via `model.save()`  
782/782 24s 31ms/step - accuracy: 0.5339 - loss: 1.30  
Epoch 5/80  
782/782 0s 30ms/step - accuracy: 0.5628 - loss: 1.23  
Epoch 5: val_accuracy improved from 0.56320 to 0.62140, saving model to best_model_separable.h5  
WARNING:absl:You are saving your model as an HDF5 file via `model.save()`  
782/782 24s 31ms/step - accuracy: 0.5628 - loss: 1.23  
Epoch 6/80  
781/782 0s 30ms/step - accuracy: 0.5988 - loss: 1.13  
Epoch 6: val_accuracy improved from 0.62140 to 0.64920, saving model to best_model_separable.h5  
WARNING:absl:You are saving your model as an HDF5 file via `model.save()`  
782/782 24s 31ms/step - accuracy: 0.5988 - loss: 1.13  
Epoch 7/80  
781/782 0s 30ms/step - accuracy: 0.6262 - loss: 1.07  
Epoch 7: val_accuracy improved from 0.64920 to 0.69960, saving model to best_model_separable.h5  
WARNING:absl:You are saving your model as an HDF5 file via `model.save()`
```

```
782/782 25s 32ms/step - accuracy: 0.6262 - loss: 1.0
Epoch 8/80
782/782 0s 29ms/step - accuracy: 0.6474 - loss: 1.01
Epoch 8: val_accuracy did not improve from 0.69960
782/782 24s 31ms/step - accuracy: 0.6474 - loss: 1.0
Epoch 9/80
782/782 0s 30ms/step - accuracy: 0.6549 - loss: 0.99
Epoch 9: val_accuracy did not improve from 0.69960
782/782 24s 31ms/step - accuracy: 0.6549 - loss: 0.9
Epoch 10/80
782/782 0s 30ms/step - accuracy: 0.6751 - loss: 0.95
Epoch 10: val_accuracy improved from 0.69960 to 0.71820, saving model to /tmp/tf/mnist/separable/2020-09-25-13-52-01/000000000000.hdf5
WARNING:absl:You are saving your model as an HDF5 file via `model.save()`
782/782 24s 31ms/step - accuracy: 0.6751 - loss: 0.9
Epoch 11/80
782/782 0s 29ms/step - accuracy: 0.6850 - loss: 0.92
Epoch 11: val_accuracy improved from 0.71820 to 0.71990, saving model to /tmp/tf/mnist/separable/2020-09-25-13-52-01/000000000001.hdf5
WARNING:absl:You are saving your model as an HDF5 file via `model.save()`
782/782 24s 31ms/step - accuracy: 0.6850 - loss: 0.9
Epoch 12/80
782/782 0s 29ms/step - accuracy: 0.6942 - loss: 0.89
Epoch 12: val_accuracy improved from 0.71990 to 0.72770, saving model to /tmp/tf/mnist/separable/2020-09-25-13-52-01/000000000002.hdf5
WARNING:absl:You are saving your model as an HDF5 file via `model.save()`
782/782 24s 31ms/step - accuracy: 0.6942 - loss: 0.8
Epoch 13/80
782/782 0s 29ms/step - accuracy: 0.7020 - loss: 0.87
Epoch 13: val_accuracy improved from 0.72770 to 0.74420, saving model to /tmp/tf/mnist/separable/2020-09-25-13-52-01/000000000003.hdf5
WARNING:absl:You are saving your model as an HDF5 file via `model.save()`
--
```

Double-click (or enter) to edit

```
test_loss, test_acc = model_separable.evaluate(x_test_norm, y_test_cat, verbose=0)
print(f" Test Accuracy: {test_acc:.4f}")
print(f" Test Loss: {test_loss:.4f}")
```

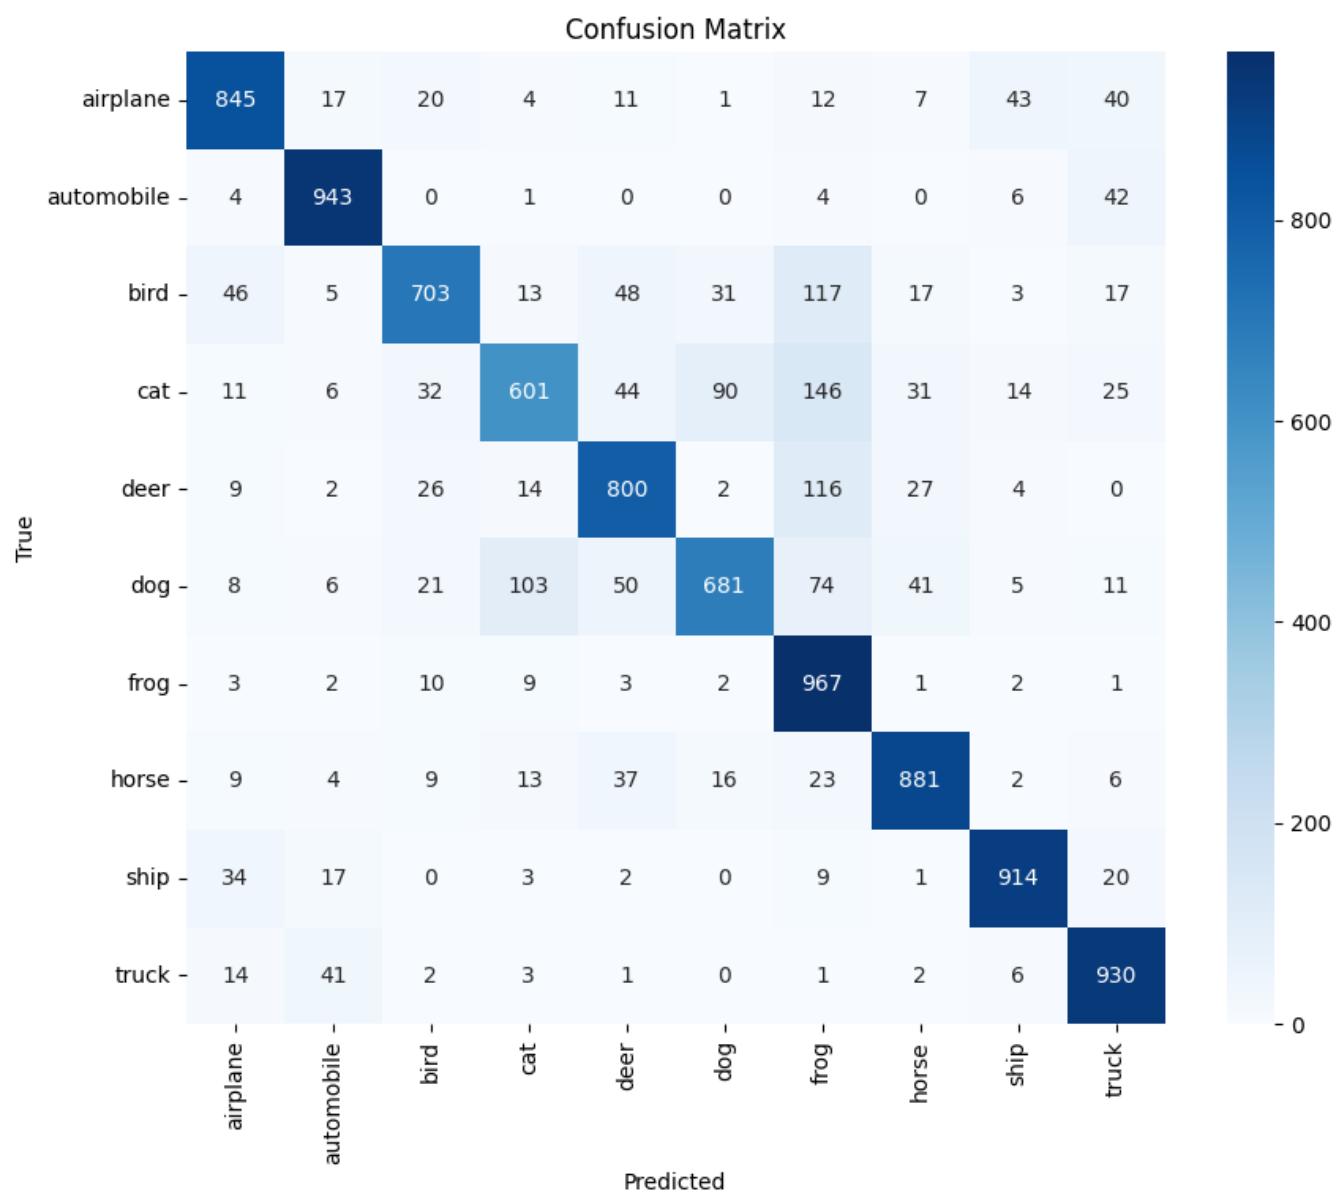
→ Test Accuracy: 0.8265
Test Loss: 0.5201

```
y_pred_probs = model_separable.predict(x_test_norm, verbose=0)
y_pred = np.argmax(y_pred_probs, axis=1)
y_true = np.argmax(y_test_cat, axis=1)
```

✓ Confusion Matrix + Heatmap

```
cm = confusion_matrix(y_true, y_pred)
plt.figure(figsize=(10, 8))
```

```
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=class_names, yti  
plt.xlabel('Predicted')  
plt.ylabel('True')  
plt.title('Confusion Matrix')  
plt.show()
```



▼ Classification Report

```
report = classification_report(y_true, y_pred, target_names=class_names)
print(" Classification Report model 2:\n")
print(report)
```

→ Classification Report model 2:

	precision	recall	f1-score	support
airplane	0.86	0.84	0.85	1000
automobile	0.90	0.94	0.92	1000
bird	0.85	0.70	0.77	1000
cat	0.79	0.60	0.68	1000
deer	0.80	0.80	0.80	1000
dog	0.83	0.68	0.75	1000
frog	0.66	0.97	0.78	1000
horse	0.87	0.88	0.88	1000
ship	0.91	0.91	0.91	1000
truck	0.85	0.93	0.89	1000
accuracy			0.83	10000
macro avg	0.83	0.83	0.82	10000
weighted avg	0.83	0.83	0.82	10000

▼ Worst case predictions

```
confidences = np.max(y_pred_probs, axis=1)
errors = np.where(y_pred != y_true)[0]
sorted_errors = errors[np.argsort(-confidences[errors])]

print(f" Worst confident wrong predictions Top 5:")
for i in sorted_errors[:5]:
    print(f"Image {i} | Predicted: {class_names[y_pred[i]]} | True: {class_name
```

→ Worst confident wrong predictions Top 5:

Image 2405 | Predicted: frog | True: cat | Confidence: 1.0000
Image 5511 | Predicted: frog | True: cat | Confidence: 0.9999
Image 3560 | Predicted: truck | True: automobile | Confidence: 0.9999
Image 2226 | Predicted: frog | True: bird | Confidence: 0.9999
Image 2495 | Predicted: ship | True: truck | Confidence: 0.9999

✓ Strongest correct predictions

```
correct = np.where(y_pred == y_true)[0]
sorted_correct = correct[np.argsort(-confidences[correct])]

print(f"\n Top confident correct predictions Top 5:")
for i in sorted_correct[:5]:
    print(f"Image {i} | Class: {class_names[y_true[i]]} | Confidence: {confider}
```



```
Top confident correct predictions Top 5:
Image 7782 | Class: ship | Confidence: 1.0000
Image 6279 | Class: automobile | Confidence: 1.0000
Image 7876 | Class: automobile | Confidence: 1.0000
Image 8975 | Class: ship | Confidence: 1.0000
Image 2947 | Class: automobile | Confidence: 1.0000
```

✓ Transfer Learning (VGG16)

```
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.applications import VGG16
from tensorflow.keras.applications.vgg16 import preprocess_input
from tensorflow.keras.models import Model
from tensorflow.keras.layers import GlobalAveragePooling2D, Dense, Dropout
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau
from tensorflow.keras.preprocessing.image import ImageDataGenerator
import tensorflow as tf
from tensorflow.keras.applications.vgg16 import VGG16, preprocess_input
```

```
y_train_cat = to_categorical(y_train, num_classes=10)
y_test_cat = to_categorical(y_test, num_classes=10)
```

```
import tensorflow as tf
from tensorflow.keras.applications import VGG16
# Import preprocess_input from the correct location
from tensorflow.keras.applications.vgg16 import preprocess_input
from tensorflow.keras.models import Model
from tensorflow.keras.layers import GlobalAveragePooling2D, Dense, Dropout
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau
from tensorflow.keras import mixed_precision
```

```
y_train_cat = to_categorical(y_train, num_classes=10)
y_test_cat = to_categorical(y_test, num_classes=10)

mixed_precision.set_global_policy('mixed_float16')
tf.config.optimizer.set_jit(True) # XLA Compilation

# Resize and preprocess
x_train_resized = tf.image.resize(x_train, [96, 96])
x_test_resized = tf.image.resize(x_test, [96, 96])
x_train_vgg = preprocess_input(x_train_resized.numpy())
x_test_vgg = preprocess_input(x_test_resized.numpy())

train_ds = tf.data.Dataset.from_tensor_slices((x_train_vgg, y_train_cat)).shuffle(1000).batch(128)
val_ds = tf.data.Dataset.from_tensor_slices((x_test_vgg, y_test_cat)).batch(128)

base_model = VGG16(include_top=False, weights='imagenet', input_shape=(96, 96, 3))
base_model.trainable = False

x = base_model.output
x = GlobalAveragePooling2D()(x)
x = Dense(64, activation='relu')(x)
x = Dropout(0.2)(x)
output = Dense(10, activation='softmax', dtype='float32')(x)

model = Model(inputs=base_model.input, outputs=output)

model.compile(optimizer=Adam(learning_rate=1e-4),
              loss='categorical_crossentropy',
              metrics=['accuracy'])

# Callbacks
early_stop = EarlyStopping(monitor='val_loss', patience=3, restore_best_weights=True)
reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=2, verbose=1)

# Phase 1: Train head
model.fit(train_ds, validation_data=val_ds, epochs=30,
          callbacks=[early_stop, reduce_lr], verbose=1)

# Fine-tuning:4 VGG
for layer in base_model.layers[-4:]:
    layer.trainable = True

model.compile(optimizer=Adam(learning_rate=1e-5),
```

```
        loss='categorical_crossentropy',
        metrics=['accuracy'])

# Phase 2: Fine-tune
model.fit(train_ds, validation_data=val_ds, epochs=10,
          callbacks=[early_stop, reduce_lr], verbose=1)

# Evaluation
test_loss, test_acc = model.evaluate(val_ds, verbose=0)
print(f"Test Accuracy: {test_acc:.4f}, Test Loss: {test_loss:.4f}")
```

→ Epoch 1/30
391/391  30s 45ms/step - accuracy: 0.2319 - loss: 5.2
Epoch 2/30
391/391  9s 22ms/step - accuracy: 0.5192 - loss: 1.57
Epoch 3/30
391/391  9s 22ms/step - accuracy: 0.6039 - loss: 1.20
Epoch 4/30
391/391  9s 22ms/step - accuracy: 0.6553 - loss: 1.03
Epoch 5/30
391/391  9s 23ms/step - accuracy: 0.6904 - loss: 0.91
Epoch 6/30
391/391  9s 22ms/step - accuracy: 0.7152 - loss: 0.83
Epoch 7/30
391/391  9s 22ms/step - accuracy: 0.7316 - loss: 0.79
Epoch 8/30
391/391  9s 22ms/step - accuracy: 0.7432 - loss: 0.75
Epoch 9/30
391/391  9s 22ms/step - accuracy: 0.7589 - loss: 0.71
Epoch 10/30
391/391  9s 22ms/step - accuracy: 0.7685 - loss: 0.67
Epoch 11/30
391/391  9s 22ms/step - accuracy: 0.7753 - loss: 0.65
Epoch 12/30
391/391  9s 22ms/step - accuracy: 0.7819 - loss: 0.62
Epoch 13/30
391/391  9s 22ms/step - accuracy: 0.7900 - loss: 0.60
Epoch 14/30
391/391  9s 22ms/step - accuracy: 0.7939 - loss: 0.59
Epoch 15/30
391/391  9s 22ms/step - accuracy: 0.7986 - loss: 0.58
Epoch 16/30
391/391  9s 23ms/step - accuracy: 0.8060 - loss: 0.56
Epoch 17/30
391/391  9s 22ms/step - accuracy: 0.8118 - loss: 0.54
Epoch 18/30
391/391  9s 22ms/step - accuracy: 0.8128 - loss: 0.54
Epoch 19/30
391/391  9s 22ms/step - accuracy: 0.8162 - loss: 0.52
Epoch 20/30
391/391  9s 22ms/step - accuracy: 0.8210 - loss: 0.52
Epoch 21/30

```
391/391 ━━━━━━━━━━ 9s 22ms/step - accuracy: 0.8219 - loss: 0.50
Epoch 22/30
391/391 ━━━━━━━━━━ 9s 22ms/step - accuracy: 0.8237 - loss: 0.50
Epoch 23/30
391/391 ━━━━━━━━━━ 9s 22ms/step - accuracy: 0.8300 - loss: 0.49
Epoch 24/30
391/391 ━━━━━━━━━━ 9s 22ms/step - accuracy: 0.8342 - loss: 0.47
Epoch 25/30
391/391 ━━━━━━━━━━ 9s 22ms/step - accuracy: 0.8362 - loss: 0.47
Epoch 26/30
391/391 ━━━━━━━━━━ 9s 22ms/step - accuracy: 0.8335 - loss: 0.47
Epoch 27/30
391/391 ━━━━━━━━━━ 9s 22ms/step - accuracy: 0.8376 - loss: 0.46
Epoch 28/30
391/391 ━━━━━━━━━━ 9s 22ms/step - accuracy: 0.8405 - loss: 0.45
Epoch 29/30
391/391 ━━━━━━━━━━ 9s 22ms/step - accuracy: 0.8418 - loss: 0.45
```

```
import tensorflow as tf
from tensorflow.keras.applications import VGG16
from tensorflow.keras.applications.vgg16 import preprocess_input
from tensorflow.keras.models import Model
from tensorflow.keras.layers import GlobalAveragePooling2D, Dense, Dropout
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau
from tensorflow.keras import mixed_precision
from tensorflow.keras.utils import to_categorical
import matplotlib.pyplot as plt

# One-hot encoding
y_train_cat = to_categorical(y_train, num_classes=10)
y_test_cat = to_categorical(y_test, num_classes=10)

#faster float16
mixed_precision.set_global_policy('mixed_float16')
tf.config.optimizer.set_jit(True) # XLA Compilation

# Resize and preprocess
x_train_resized = tf.image.resize(x_train, [96, 96])
x_test_resized = tf.image.resize(x_test, [96, 96])
x_train_vgg = preprocess_input(x_train_resized.numpy())
x_test_vgg = preprocess_input(x_test_resized.numpy())

#
train_ds = tf.data.Dataset.from_tensor_slices((x_train_vgg, y_train_cat)).shuffle(1000).batch(128)
val_ds = tf.data.Dataset.from_tensor_slices((x_test_vgg, y_test_cat)).batch(128)

# تحميل VGG16
base_model = VGG16(include_top=False, weights='imagenet', input_shape=(96, 96, 3))
```

```
base_model.trainable = False

# رأس النموذج
x = base_model.output
x = GlobalAveragePooling2D()(x)
x = Dense(64, activation='relu')(x)
x = Dropout(0.2)(x)
output = Dense(10, activation='softmax', dtype='float32')(x) # لازم float32 فی الم

model = Model(inputs=base_model.input, outputs=output)

model.compile(optimizer=Adam(learning_rate=1e-4),
              loss='categorical_crossentropy',
              metrics=['accuracy'])

# Callbacks
early_stop = EarlyStopping(monitor='val_loss', patience=3, restore_best_weights=True)
reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=2, verbose=1)

# ===== Phase 1: Train head =====
history1 = model.fit(train_ds, validation_data=val_ds, epochs=80,
                     callbacks=[early_stop, reduce_lr], verbose=1)

# ===== Phase 2: Fine-tuning =====
for layer in base_model.layers[-4:]:
    layer.trainable = True

model.compile(optimizer=Adam(learning_rate=1e-5),
              loss='categorical_crossentropy',
              metrics=['accuracy'])

history2 = model.fit(train_ds, validation_data=val_ds, epochs=10,
                     callbacks=[early_stop, reduce_lr], verbose=1)
```

→ Epoch 1/30
391/391 16s 32ms/step - accuracy: 0.2228 - loss: 6.266
Epoch 2/30
391/391 9s 23ms/step - accuracy: 0.5066 - loss: 1.6483
Epoch 3/30
391/391 9s 23ms/step - accuracy: 0.6035 - loss: 1.2258
Epoch 4/30
391/391 9s 22ms/step - accuracy: 0.6568 - loss: 1.0403
Epoch 5/30
391/391 9s 23ms/step - accuracy: 0.6949 - loss: 0.9248
Epoch 6/30
391/391 9s 22ms/step - accuracy: 0.7183 - loss: 0.8461
Epoch 7/30
391/391 9s 22ms/step - accuracy: 0.7375 - loss: 0.7901

Epoch 8/30
391/391 9s 23ms/step - accuracy: 0.7461 - loss: 0.7483

Epoch 9/30
391/391 9s 23ms/step - accuracy: 0.7643 - loss: 0.7053

Epoch 10/30
391/391 9s 22ms/step - accuracy: 0.7697 - loss: 0.6819

Epoch 11/30
391/391 9s 23ms/step - accuracy: 0.7802 - loss: 0.6497

Epoch 12/30
391/391 9s 22ms/step - accuracy: 0.7883 - loss: 0.6265

Epoch 13/30
391/391 9s 22ms/step - accuracy: 0.7938 - loss: 0.6072

Epoch 14/30
391/391 9s 22ms/step - accuracy: 0.7986 - loss: 0.5941

Epoch 15/30
391/391 9s 22ms/step - accuracy: 0.8042 - loss: 0.5770

Epoch 16/30
391/391 9s 22ms/step - accuracy: 0.8059 - loss: 0.5614

Epoch 17/30
391/391 9s 23ms/step - accuracy: 0.8141 - loss: 0.5491

Epoch 18/30
391/391 9s 22ms/step - accuracy: 0.8145 - loss: 0.5395

Epoch 19/30
391/391 9s 22ms/step - accuracy: 0.8206 - loss: 0.5271

Epoch 20/30
391/391 9s 22ms/step - accuracy: 0.8225 - loss: 0.5185

Epoch 21/30
391/391 9s 22ms/step - accuracy: 0.8262 - loss: 0.5070

Epoch 22/30
391/391 9s 22ms/step - accuracy: 0.8281 - loss: 0.4996

Epoch 23/30
391/391 9s 22ms/step - accuracy: 0.8316 - loss: 0.4881

Epoch 24/30
391/391 9s 22ms/step - accuracy: 0.8331 - loss: 0.4858

Epoch 25/30
391/391 9s 22ms/step - accuracy: 0.8360 - loss: 0.4803

Epoch 26/30
391/391 9s 22ms/step - accuracy: 0.8398 - loss: 0.4666

Epoch 27/30
391/391 9s 22ms/step - accuracy: 0.8426 - loss: 0.4598

Epoch 28/30
391/391 9s 23ms/step - accuracy: 0.8450 - loss: 0.4547

Epoch 29/30
391/391 9s 23ms/step - accuracy: 0.8429 - loss: 0.4522

Epoch 30/30
391/391 9s 22ms/step - accuracy: 0.8477 - loss: 0.4432

Restoring model weights from the end of the best epoch: 29.

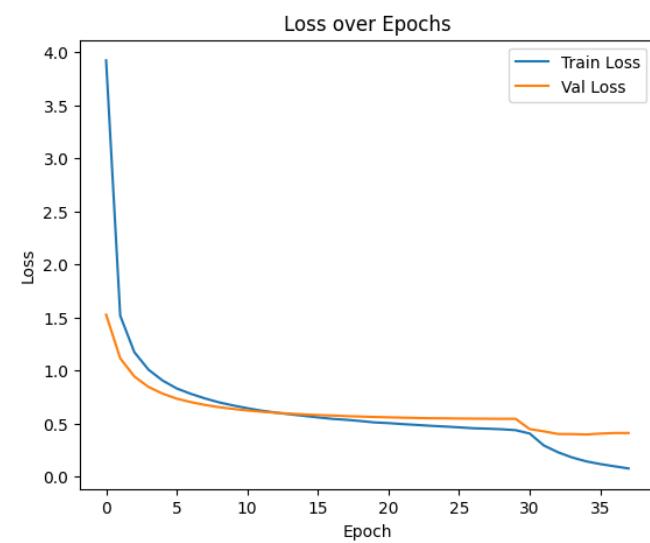
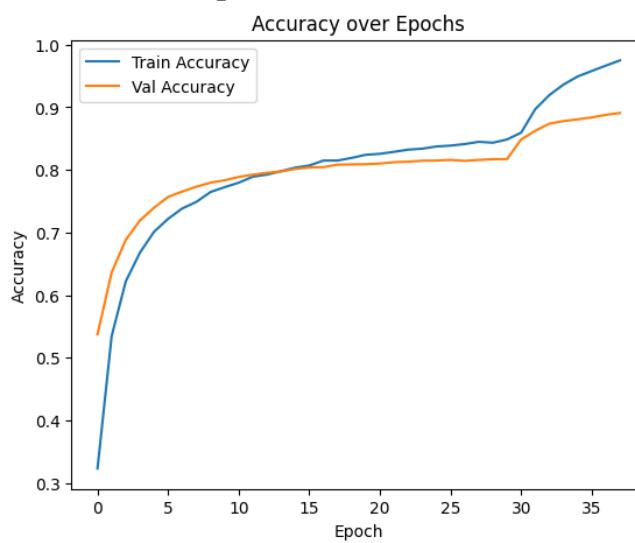
Epoch 1/10
391/391 18s 36ms/step - accuracy: 0.8527 - loss: 0.426

Epoch 2/10
391/391 10s 26ms/step - accuracy: 0.8959 - loss: 0.299

Epoch 3/10
391/391 10s 26ms/step - accuracy: 0.9168 - loss: 0.236

Epoch 4/10

```
391/391 ━━━━━━━━━━ 10s 26ms/step - accuracy: 0.9341 - loss: 0.187
Epoch 5/10
391/391 ━━━━━━━━ 10s 26ms/step - accuracy: 0.9471 - loss: 0.149
Epoch 6/10
391/391 ━━━━━━ 10s 26ms/step - accuracy: 0.9569 - loss: 0.121
Epoch 7/10
391/391 ━━━━ 0s 22ms/step - accuracy: 0.9656 - loss: 0.1017
Epoch 7: ReduceLROnPlateau reducing learning rate to 4.999999873689376e-06.
391/391 ━━━━ 10s 26ms/step - accuracy: 0.9656 - loss: 0.101
Epoch 8/10
391/391 ━━━━ 10s 26ms/step - accuracy: 0.9735 - loss: 0.081
Epoch 8: early stopping
Restoring model weights from the end of the best epoch: 5.
Test Accuracy: 0.8809, Test Loss: 0.3976
```



```
for layer in model.layers:  
    print(layer.name)
```

```
→ input_layer  
block1_conv1  
block1_conv2  
block1_pool  
block2_conv1  
block2_conv2  
block2_pool  
block3_conv1  
block3_conv2  
block3_conv3  
block3_pool  
block4_conv1  
block4_conv2  
block4_conv3  
block4_pool  
block5_conv1  
block5_conv2  
block5_conv3  
block5_pool  
global_average_pooling2d  
dense  
dropout  
dense_1
```

```
model.save("vgg_light_finetuned22.h5")
```

```
→ WARNING:absl:You are saving your model as an HDF5 file via `model.save()` o
```

✓ model summary

```
model.summary()
```

```
→ Model: "functional"
```

Layer (type)	Output Shape	Param #
input_layer (InputLayer)	(None, 96, 96, 3)	0
cast (Cast)	(None, 96, 96, 3)	0
block1_conv1 (Conv2D)	(None, 96, 96, 64)	1,792
block1_conv2 (Conv2D)	(None, 96, 96, 64)	36,928
block1_pool (MaxPooling2D)	(None, 48, 48, 64)	0

block2_conv1 (Conv2D)	(None, 48, 48, 128)	73,856
block2_conv2 (Conv2D)	(None, 48, 48, 128)	147,584
block2_pool (MaxPooling2D)	(None, 24, 24, 128)	0
block3_conv1 (Conv2D)	(None, 24, 24, 256)	295,168
block3_conv2 (Conv2D)	(None, 24, 24, 256)	590,080
block3_conv3 (Conv2D)	(None, 24, 24, 256)	590,080
block3_pool (MaxPooling2D)	(None, 12, 12, 256)	0
block4_conv1 (Conv2D)	(None, 12, 12, 512)	1,180,160
block4_conv2 (Conv2D)	(None, 12, 12, 512)	2,359,808
block4_conv3 (Conv2D)	(None, 12, 12, 512)	2,359,808
block4_pool (MaxPooling2D)	(None, 6, 6, 512)	0
block5_conv1 (Conv2D)	(None, 6, 6, 512)	2,359,808
block5_conv2 (Conv2D)	(None, 6, 6, 512)	2,359,808
block5_conv3 (Conv2D)	(None, 6, 6, 512)	2,359,808
block5_pool (MaxPooling2D)	(None, 3, 3, 512)	0
global_average_pooling2d (GlobalAveragePooling2D)	(None, 512)	0
dense (Dense)	(None, 64)	32,832
dropout (Dropout)	(None, 64)	0
cast_1 (Cast)	(None, 64)	0
dense_1 (Dense)	(None, 10)	650

Total params: 28,973,988 (110.53 MB)

Trainable params: 7,112,906 (27.13 MB)

Non-trainable params: 7,635,264 (29.13 MB)

Optimizer params: 14,225,818 (54.27 MR)

✓ Evaluate the model on test set

```
from sklearn.metrics import classification_report, confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np

# توقعات النموذج
y_pred_probs = model.predict(val_ds)
y_pred = np.argmax(y_pred_probs, axis=1)
y_true = np.argmax(y_test_cat, axis=1)

# Classification Report
print("Classification Report:\n", classification_report(y_true, y_pred))
```

79/79 ————— 3s 29ms/step

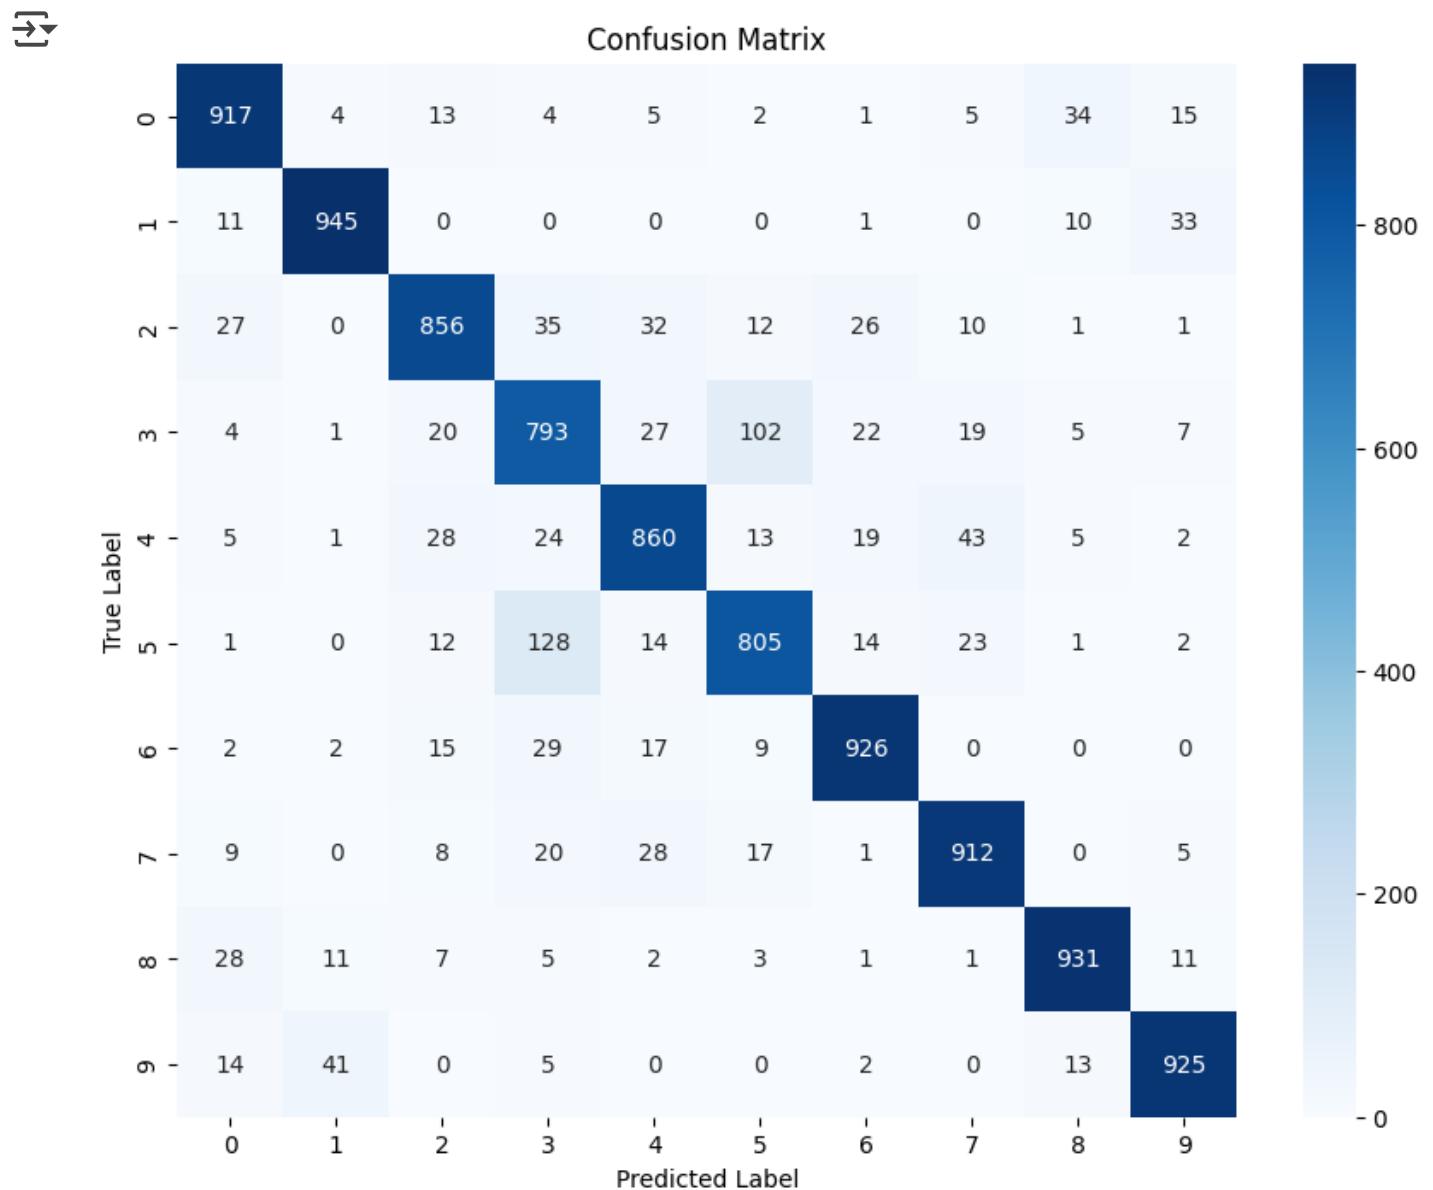
Classification Report:

	precision	recall	f1-score	support
0	0.90	0.92	0.91	1000
1	0.94	0.94	0.94	1000
2	0.89	0.86	0.87	1000
3	0.76	0.79	0.78	1000
4	0.87	0.86	0.87	1000
5	0.84	0.81	0.82	1000
6	0.91	0.93	0.92	1000
7	0.90	0.91	0.91	1000
8	0.93	0.93	0.93	1000
9	0.92	0.93	0.92	1000
accuracy			0.89	10000
macro avg	0.89	0.89	0.89	10000
weighted avg	0.89	0.89	0.89	10000

Start coding or generate with AI.

▼ Confusion Matrix

```
conf_mat = confusion_matrix(y_true, y_pred)
plt.figure(figsize=(10, 8))
sns.heatmap(conf_mat, annot=True, fmt='d', cmap='Blues')
plt.title('Confusion Matrix')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```



✓ High Confidence Errors

```
wrong_indices = np.where(y_pred != y_true)[0]
wrong_confidences = np.max(y_pred_probs[wrong_indices], axis=1)
top_wrong = wrong_indices[np.argsort(wrong_confidences)[-10:]]

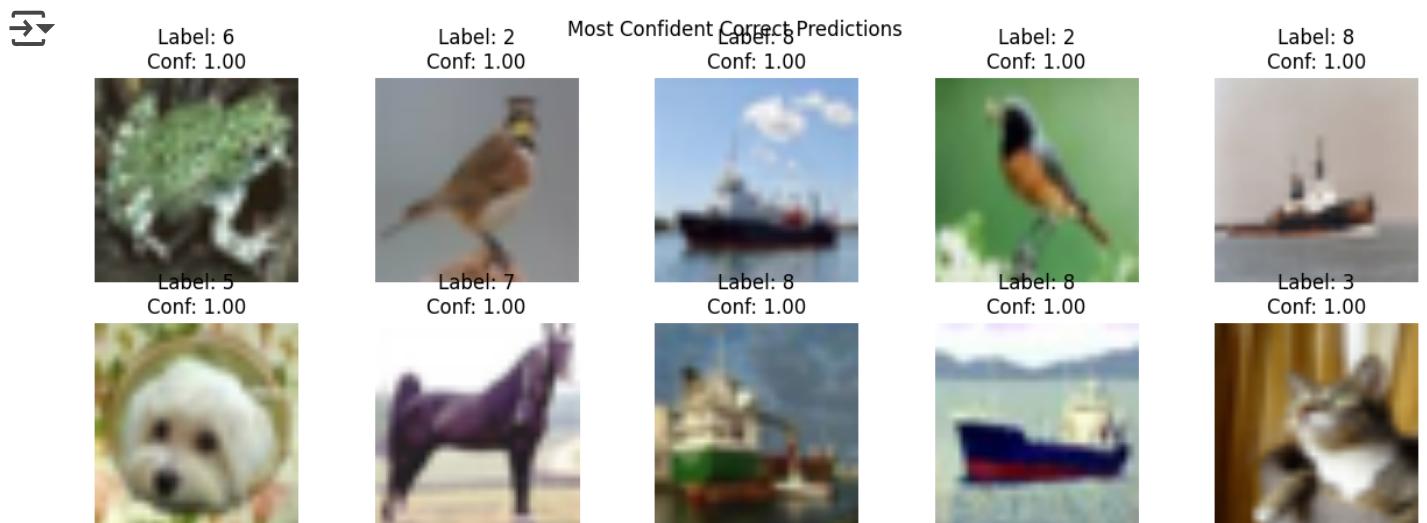
plt.figure(figsize=(15, 5))
for i, idx in enumerate(top_wrong):
    plt.subplot(2, 5, i + 1)
    plt.imshow(x_test_resized[idx].numpy().astype("uint8"))
    plt.title(f"True: {y_true[idx]}, Pred: {y_pred[idx]}\nConf: {wrong_confidences[idx]}")
    plt.axis("off")
plt.suptitle("Most Confident Wrong Predictions")
plt.show()
```



✓ Top Confident Correct Predictions

```
correct_indices = np.where(y_pred == y_true)[0]
correct_confidences = np.max(y_pred_probs[correct_indices], axis=1)
top_correct = correct_indices[np.argsort(correct_confidences)[-10:]]

plt.figure(figsize=(15, 5))
for i, idx in enumerate(top_correct):
    plt.subplot(2, 5, i + 1)
    plt.imshow(x_test_resized[idx].numpy().astype("uint8"))
    plt.title(f"Label: {y_true[idx]}\nConf: {correct_confidences[np.argsort(correct_confidences)[-10+idx]]}")
    plt.axis("off")
plt.suptitle("Most Confident Correct Predictions")
plt.show()
```



✓ Accuracy & Loss Curves

```
test_loss, test_acc = model.evaluate(val_ds, verbose=0)
```

```
print(f"Test Accuracy: {test_acc:.4f}, Test Loss: {test_loss:.4f}")

def merge_histories(h1, h2):
    merged = {}
    for key in h1.history:
        merged[key] = h1.history[key] + h2.history[key]
    return merged

def plot_training_curves(history):
    acc = history['accuracy']
    val_acc = history['val_accuracy']
    loss = history['loss']
    val_loss = history['val_loss']

    plt.figure(figsize=(14, 5))

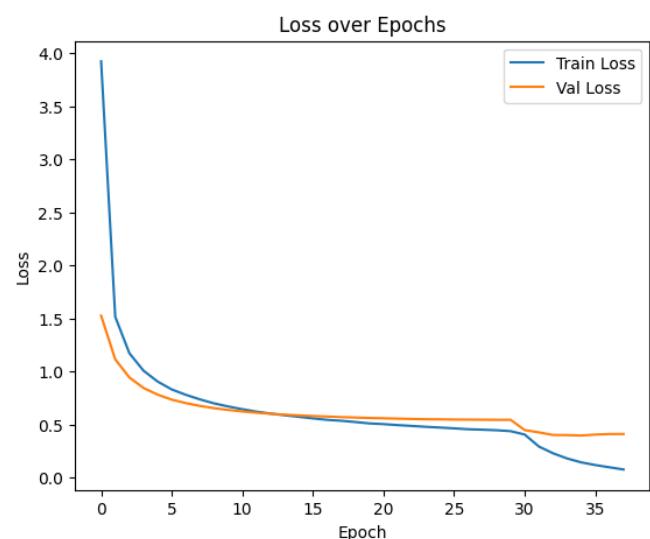
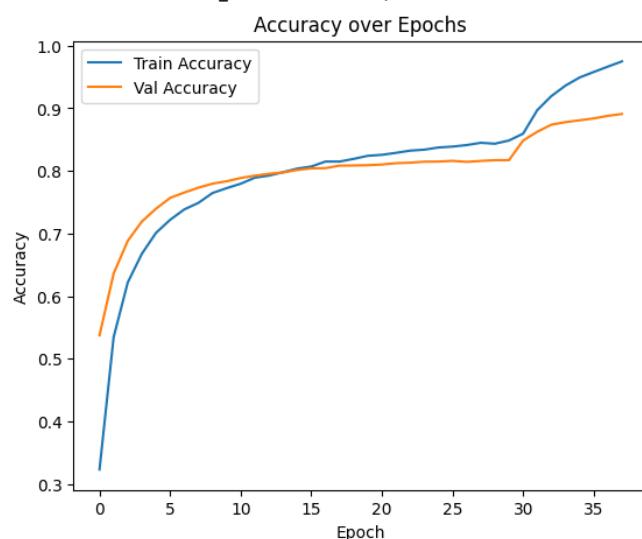
    # Accuracy
    plt.subplot(1, 2, 1)
    plt.plot(acc, label='Train Accuracy')
    plt.plot(val_acc, label='Val Accuracy')
    plt.title('Accuracy over Epochs')
    plt.xlabel('Epoch')
    plt.ylabel('Accuracy')
    plt.legend()

    # Loss
    plt.subplot(1, 2, 2)
    plt.plot(loss, label='Train Loss')
    plt.plot(val_loss, label='Val Loss')
    plt.title('Loss over Epochs')
    plt.xlabel('Epoch')
    plt.ylabel('Loss')
    plt.legend()

    plt.show()

merged_history = merge_histories(history1, history2)
plot_training_curves(merged_history)
```

→ Test Accuracy: 0.8809, Test Loss: 0.3976



✓ Grad-CAM on VGG model

```
import numpy as np
import tensorflow as tf
import matplotlib.pyplot as plt
import cv2

index = 7
img_input = x_test_vgg[index]
original_img = x_test_resized[index].numpy()

img_tensor = tf.convert_to_tensor(np.expand_dims(img_input, axis=0))

last_conv_layer_name = "block5_conv3"
```

```
grad_model = tf.keras.models.Model(  
    [model.inputs],  
    [model.get_layer(last_conv_layer_name).output, model.output]  
)  
  
with tf.GradientTape() as tape:  
    conv_outputs, predictions = grad_model(img_tensor)  
    predicted_class = tf.argmax(predictions[0])  
    loss = predictions[:, predicted_class]  
  
    grads = tape.gradient(loss, conv_outputs)  
    pooled_grads = tf.reduce_mean(grads, axis=(0, 1, 2))  
  
    conv_outputs = conv_outputs[0]  
    heatmap = conv_outputs @ pooled_grads[..., tf.newaxis]  
    heatmap = tf.squeeze(heatmap)  
  
    heatmap = tf.squeeze(heatmap)  
  
if len(heatmap.shape) != 2 or tf.reduce_max(heatmap).numpy() == 0:  
    heatmap = tf.ones((12, 12))  
  
    heatmap = np.maximum(heatmap, 0)  
    heatmap = heatmap / np.max(heatmap)  
  
    heatmap = np.uint8(255 * heatmap)  
    heatmap = cv2.resize(heatmap, (96, 96))  
    heatmap_color = cv2.applyColorMap(heatmap, cv2.COLORMAP_JET)  
  
    original_img_norm = (original_img - original_img.min()) / (original_img.max() -  
    original_img_uint8 = np.uint8(original_img_norm * 255)  
  
    alpha = 0.4  
    overlay = cv2.addWeighted(original_img_uint8, 1 - alpha, heatmap_color, alpha,  
  
    plt.figure(figsize=(12, 4))  
  
    plt.subplot(1, 3, 1)
```

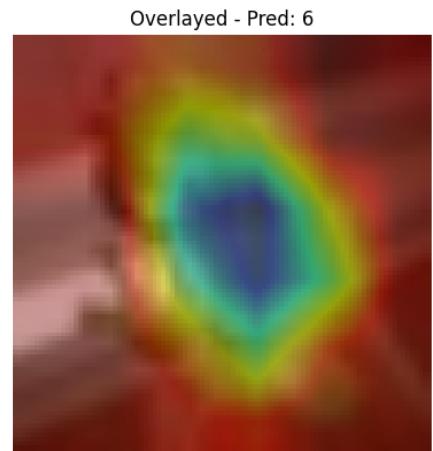
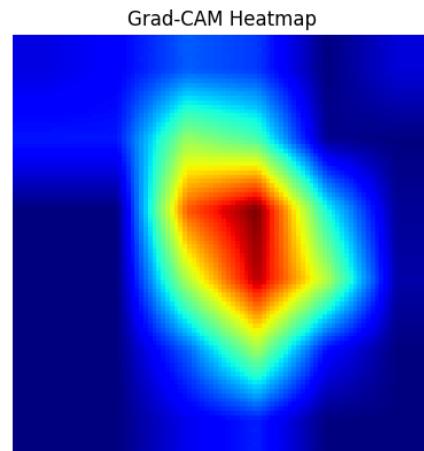
```
plt.title("Original Image")
plt.imshow(original_img_uint8.astype('uint8'))
plt.axis('off')

plt.subplot(1, 3, 2)
plt.title("Grad-CAM Heatmap")
plt.imshow(heatmap, cmap='jet')
plt.axis('off')

plt.subplot(1, 3, 3)
plt.title(f"Overlaid - Pred: {predicted_class.numpy()}")
plt.imshow(overlay)
plt.axis('off')

plt.tight_layout()
plt.show()
```

→ /usr/local/lib/python3.11/dist-packages/keras/src/models/functional.py:237:
Expected: [['keras_tensor_25']]
Received: inputs=Tensor(shape=(1, 96, 96, 3))
raise ValueError(msg)



✓ The Bonus Autoencoder + KMeans + visualization

```
!pip install tensorflow --quiet

import numpy as np
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
from sklearn.manifold import TSNE
from sklearn.metrics import accuracy_score, adjusted_rand_score
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Input, Conv2D, MaxPooling2D, UpSampling2D
from tensorflow.keras.datasets import cifar10
from tensorflow.keras.optimizers import Adam
import seaborn as sns

# ===== Step 1: Load and normalize CIFAR-10 =====
(x_train, y_train), (x_test, y_test) = cifar10.load_data()
x_data = np.concatenate((x_train, x_test), axis=0).astype('float32') / 255.0
y_data = np.concatenate((y_train, y_test), axis=0).flatten()

# ===== Step 2: Build Autoencoder =====
input_img = Input(shape=(32, 32, 3))
x = Conv2D(32, 3, activation='relu', padding='same')(input_img)
x = MaxPooling2D(2, padding='same')(x)
x = Conv2D(64, 3, activation='relu', padding='same')(x)
encoded = MaxPooling2D(2, padding='same')(x) # (8, 8, 64)

x = Conv2D(64, 3, activation='relu', padding='same')(encoded)
x = UpSampling2D(2)(x)
x = Conv2D(32, 3, activation='relu', padding='same')(x)
x = UpSampling2D(2)(x)
decoded = Conv2D(3, 3, activation='sigmoid', padding='same')(x)

autoencoder = Model(input_img, decoded)
autoencoder.compile(optimizer=Adam(1e-3), loss='mse')

# ===== Step 3: Train Autoencoder =====
autoencoder.fit(x_data[:40000], x_data[:40000],
                 epochs=15, batch_size=256,
                 validation_split=0.1, shuffle=True)

# ===== Step 4: Encode Data =====
encoder = Model(input_img, encoded)
features = encoder.predict(x_data[:10000])
```

```
features_flat = features.reshape((features.shape[0], -1))
true_labels = y_data[:10000]

# ===== Step 5: Clustering with KMeans =====
kmeans = KMeans(n_clusters=10, random_state=42)
pred_clusters = kmeans.fit_predict(features_flat)

# ===== Step 6: Evaluate clustering (Unsupervised Accuracy) =====
# ملاحظة: في التعلم غير الخاضع للإشراف لا يوجد # كمقياس رسمي
ari_score = adjusted_rand_score(true_labels, pred_clusters)
print(f"Adjusted Rand Index (ARI): {ari_score:.4f}")

# ===== Step 7: Visualize Clusters with t-SNE =====
tsne = TSNE(n_components=2, random_state=42)
tsne_result = tsne.fit_transform(features_flat)

plt.figure(figsize=(10, 6))
sns.scatterplot(x=tsne_result[:, 0], y=tsne_result[:, 1], hue=pred_clusters, palette='viridis')
plt.title("t-SNE of Encoded Features + KMeans Clustering")
plt.xlabel("t-SNE Component 1")
plt.ylabel("t-SNE Component 2")
plt.grid(True)
plt.legend(title='Cluster')
plt.show()

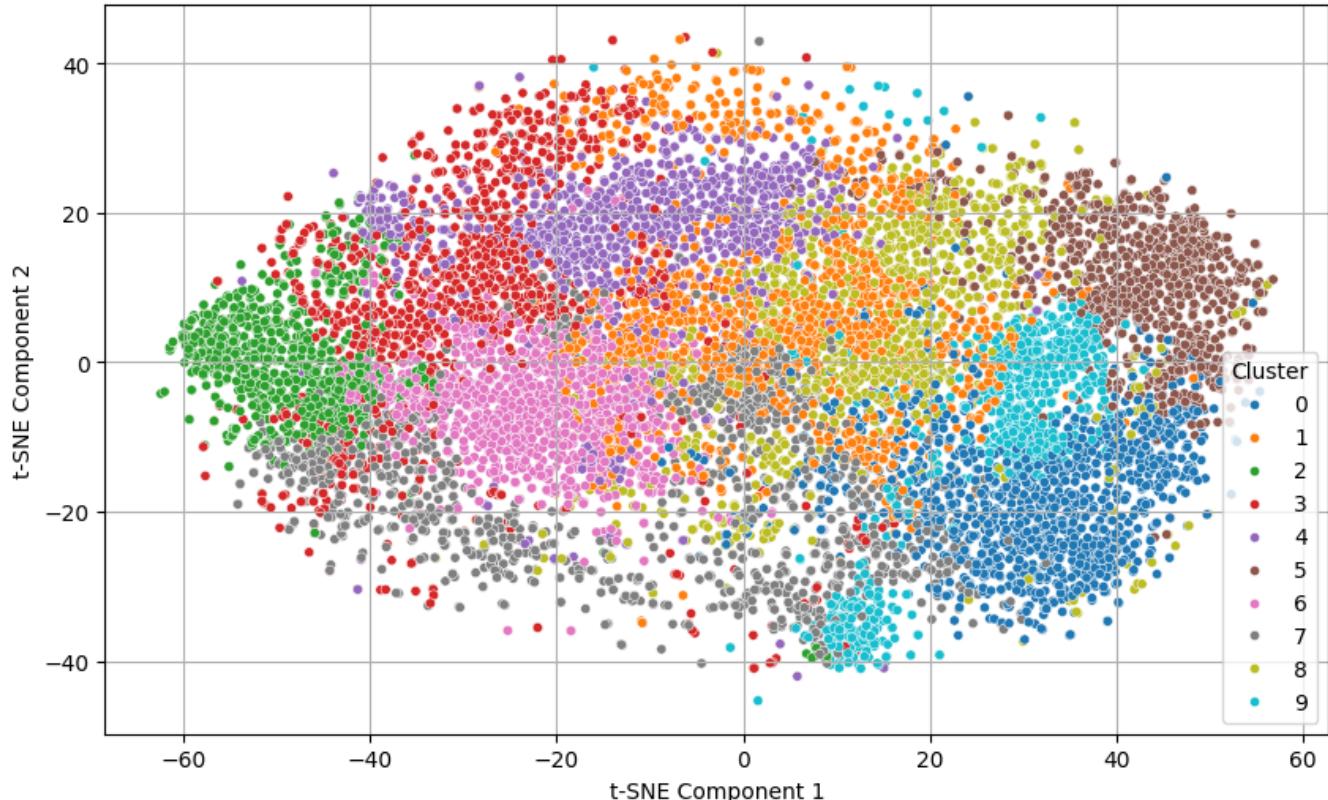
# ===== Step 8: Show example images from each cluster =====
def show_cluster_examples(images, labels, cluster_ids, samples_per_cluster=5):
    plt.figure(figsize=(samples_per_cluster*2, 20))
    for cluster in range(10):
        idxs = np.where(cluster_ids == cluster)[0][:samples_per_cluster]
        for i, idx in enumerate(idxs):
            plt.subplot(10, samples_per_cluster, cluster*samples_per_cluster + i + 1)
            plt.imshow(images[idx])
            plt.axis('off')
        if i == 0:
            plt.ylabel(f"Cluster {cluster}", fontsize=12)
    plt.suptitle("Example Images from Each Cluster", fontsize=16)
    plt.tight_layout()
    plt.show()

show_cluster_examples(x_data[:10000], true_labels, pred_clusters)
```

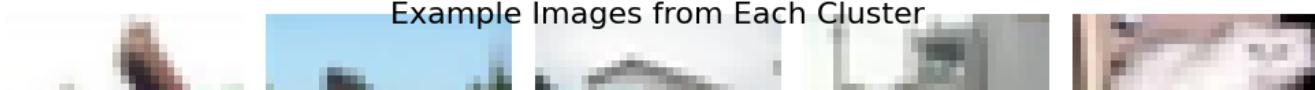
→ Epoch 1/15
141/141 ━━━━━━━━━━ 7s 26ms/step - loss: 0.0322 - val_loss: 0.0083
Epoch 2/15
141/141 ━━━━━━━━ 1s 6ms/step - loss: 0.0077 - val_loss: 0.0064
Epoch 3/15

```
141/141 ━━━━━━━━━━ 1s 6ms/step - loss: 0.0062 - val_loss: 0.0056
Epoch 4/15
141/141 ━━━━━━━━━━ 1s 6ms/step - loss: 0.0055 - val_loss: 0.0051
Epoch 5/15
141/141 ━━━━━━━━━━ 1s 6ms/step - loss: 0.0050 - val_loss: 0.0047
Epoch 6/15
141/141 ━━━━━━━━━━ 1s 6ms/step - loss: 0.0046 - val_loss: 0.0044
Epoch 7/15
141/141 ━━━━━━━━━━ 1s 6ms/step - loss: 0.0044 - val_loss: 0.0045
Epoch 8/15
141/141 ━━━━━━━━━━ 1s 6ms/step - loss: 0.0042 - val_loss: 0.0041
Epoch 9/15
141/141 ━━━━━━━━━━ 1s 6ms/step - loss: 0.0041 - val_loss: 0.0039
Epoch 10/15
141/141 ━━━━━━━━━━ 1s 6ms/step - loss: 0.0040 - val_loss: 0.0042
Epoch 11/15
141/141 ━━━━━━━━━━ 1s 6ms/step - loss: 0.0038 - val_loss: 0.0037
Epoch 12/15
141/141 ━━━━━━━━━━ 1s 6ms/step - loss: 0.0038 - val_loss: 0.0036
Epoch 13/15
141/141 ━━━━━━━━━━ 1s 6ms/step - loss: 0.0036 - val_loss: 0.0037
Epoch 14/15
141/141 ━━━━━━━━━━ 1s 6ms/step - loss: 0.0036 - val_loss: 0.0035
Epoch 15/15
141/141 ━━━━━━━━━━ 1s 6ms/step - loss: 0.0035 - val_loss: 0.0035
313/313 ━━━━━━━━━━ 1s 2ms/step
Adjusted Rand Index (ARI): 0.0511
```

t-SNE of Encoded Features + KMeans Clustering



Example Images from Each Cluster



```
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
from sklearn.cluster import KMeans
from sklearn.manifold import TSNE
from sklearn.metrics import adjusted_rand_score
from tensorflow.keras.datasets import cifar10
from tensorflow.keras.applications.vgg16 import VGG16, preprocess_input
from tensorflow.keras.models import Model
from tensorflow.keras.layers import GlobalAveragePooling2D
import tensorflow as tf

# Step 1: Load CIFAR-10
(x_train, y_train), (x_test, y_test) = cifar10.load_data()
x_data = np.concatenate([x_train, x_test])
y_data = np.concatenate([y_train, y_test]).flatten()

# Step 2: Use 10,000 samples for better clustering
x_subset = x_data[:10000]
y_subset = y_data[:10000]

# Resize and preprocess
x_resized = tf.image.resize(x_subset, [224, 224]).numpy()
x_preprocessed = preprocess_input(x_resized)

# Step 3: VGG16 model + GlobalAveragePooling for better feature reduction
vgg = VGG16(weights='imagenet', include_top=False, input_shape=(224, 224, 3))
x = vgg.output
x = GlobalAveragePooling2D()(x)
feature_model = Model(inputs=vgg.input, outputs=x)

# Step 4: Extract features
features = feature_model.predict(x_preprocessed, batch_size=64, verbose=1)

# Step 5: KMeans clustering
kmeans = KMeans(n_clusters=10, random_state=42)
clusters = kmeans.fit_predict(features)

# Step 6: Evaluate clustering
ari = adjusted_rand_score(y_subset, clusters)
print(f"\n🎯 Adjusted Rand Index (ARI): {ari:.4f}")

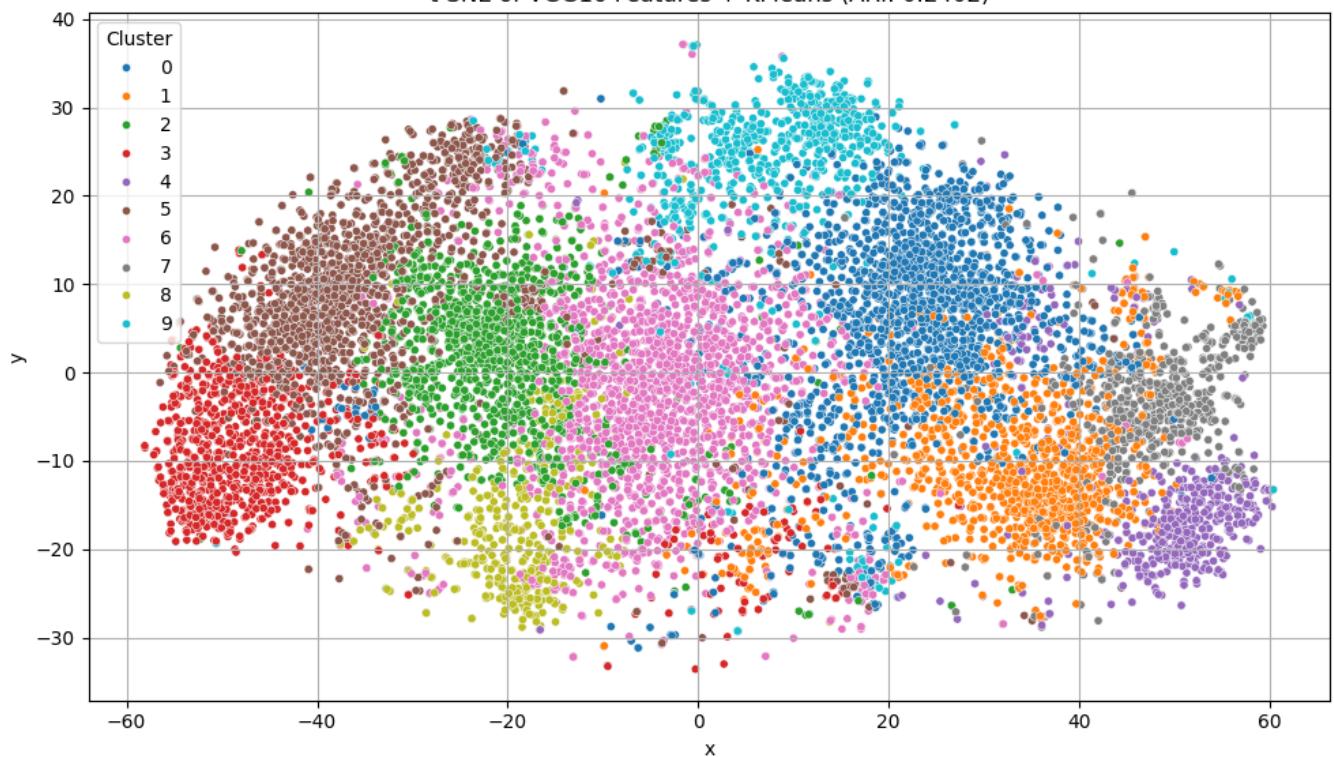
# Step 7: Visualize with t-SNE
tsne = TSNE(n_components=2, perplexity=50, random_state=42)
tsne_result = tsne.fit_transform(features)
```

```
# Plot
df = pd.DataFrame({
    'x': tsne_result[:, 0],
    'y': tsne_result[:, 1],
    'cluster': clusters
})

plt.figure(figsize=(10, 6))
sns.scatterplot(data=df, x='x', y='y', hue='cluster', palette='tab10', s=20)
plt.title(f"t-SNE of VGG16 Features + KMeans (ARI: {ari:.4f})")
plt.grid(True)
plt.legend(title='Cluster')
plt.tight_layout()
plt.show()
```

157 / 157 ━━━━━━ 48s 181ms/step

⌚ Adjusted Rand Index (ARI): 0.2402
t-SNE of VGG16 Features + KMeans (ARI: 0.2402)



```
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
from sklearn.cluster import KMeans
from sklearn.manifold import TSNE
from sklearn.metrics import adjusted_rand_score
from tensorflow.keras.datasets import cifar10
from tensorflow.keras.applications.vgg16 import VGG16, preprocess_input
from tensorflow.keras.models import Model
from tensorflow.keras.layers import GlobalAveragePooling2D, Dense, Dropout
import tensorflow as tf
```

```
# Step 1: Load CIFAR-10 and preprocess
(x_train, y_train), (x_test, y_test) = cifar10.load_data()
x_data = np.concatenate([x_train, x_test])
y_data = np.concatenate([y_train, y_test]).flatten()

x_subset = x_data[:10000]
y_subset = y_data[:10000]

x_resized = tf.image.resize(x_subset, [224, 224]).numpy()
x_preprocessed = preprocess_input(x_resized)

# Step 2: VGG16 feature model with enhancements
base_model = VGG16(weights='imagenet', include_top=False, input_shape=(224, 224))
for layer in base_model.layers:
    layer.trainable = False

x = base_model.output
x = GlobalAveragePooling2D()(x)
x = Dense(256, activation='relu')(x)
x = Dropout(0.4)(x)
x = Dense(128, activation='relu')(x)
feature_output = Dropout(0.4)(x)

feature_model = Model(inputs=base_model.input, outputs=feature_output)

# Step 3: Feature extraction
features = feature_model.predict(x_preprocessed, batch_size=64, verbose=1)

# Step 4: Clustering
kmeans = KMeans(n_clusters=10, random_state=42)
clusters = kmeans.fit_predict(features)

# Step 5: Evaluate with ARI
ari = adjusted_rand_score(y_subset, clusters)
print(f"\n🎯 Adjusted Rand Index (ARI): {ari:.4f}")

# Step 6: Save CSV
df_results = pd.DataFrame({
    'Image_Index': np.arange(len(y_subset)),
    'True_Label': y_subset,
    'Cluster': clusters
})
df_results.to_csv("vgg_clusters.csv", index=False)

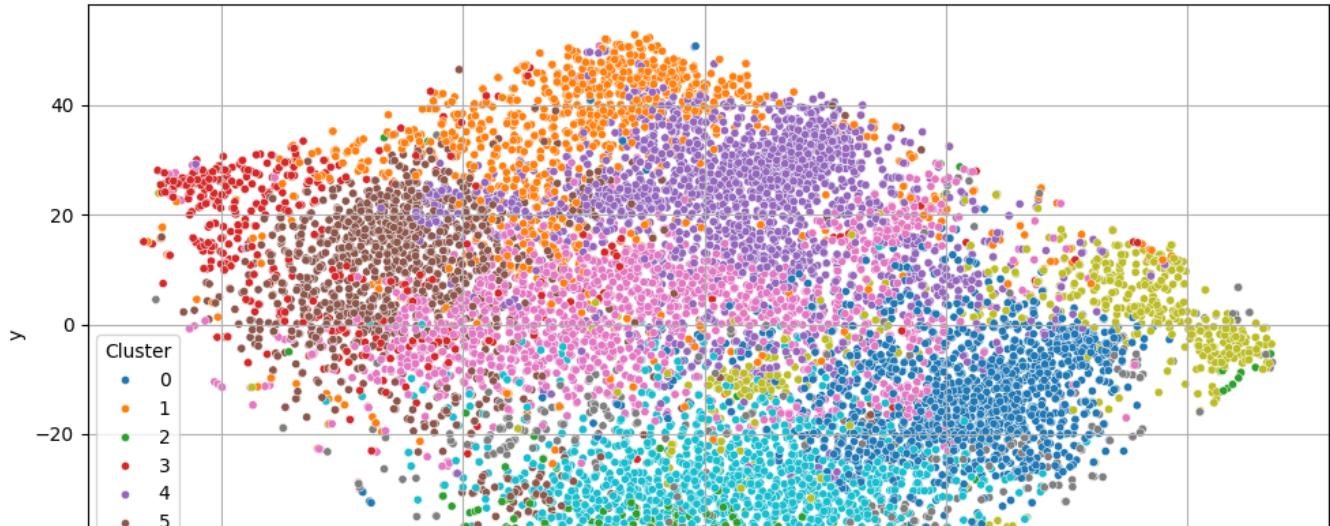
# Step 7: t-SNE plot
tsne_result = TSNE(n_components=2, perplexity=50, random_state=42).fit_transform(df_tsne = pd.DataFrame({
```

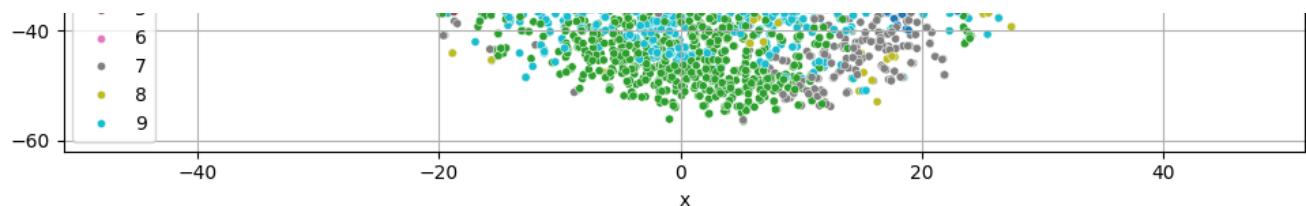
```
'x': tsne_result[:, 0],  
'y': tsne_result[:, 1],  
'cluster': clusters  
})  
  
plt.figure(figsize=(10, 6))  
sns.scatterplot(data=df_tsne, x='x', y='y', hue='cluster', palette='tab10', s=2  
plt.title(f"t-SNE of VGG16 Features + KMeans Clustering\nARI: {ari:.4f}")  
plt.grid(True)  
plt.tight_layout()  
plt.legend(title="Cluster")  
plt.show()  
  
# Step 8: Show example images from each cluster  
def show_cluster_examples(images, cluster_ids, samples_per_cluster=5):  
    plt.figure(figsize=(samples_per_cluster * 2, 20))  
    for cluster in range(10):  
        idxs = np.where(cluster_ids == cluster)[0][:samples_per_cluster]  
        for i, idx in enumerate(idxs):  
            plt.subplot(10, samples_per_cluster, cluster * samples_per_cluster + i + 1)  
            plt.imshow(images[idx])  
            plt.axis('off')  
        if i == 0:  
            plt.ylabel(f"Cluster {cluster}", fontsize=12)  
    plt.suptitle("Example Images from Each Cluster", fontsize=16)  
    plt.tight_layout()  
    plt.show()  
  
show_cluster_examples(x_subset, clusters)
```

157 / 157 25s 154ms/step

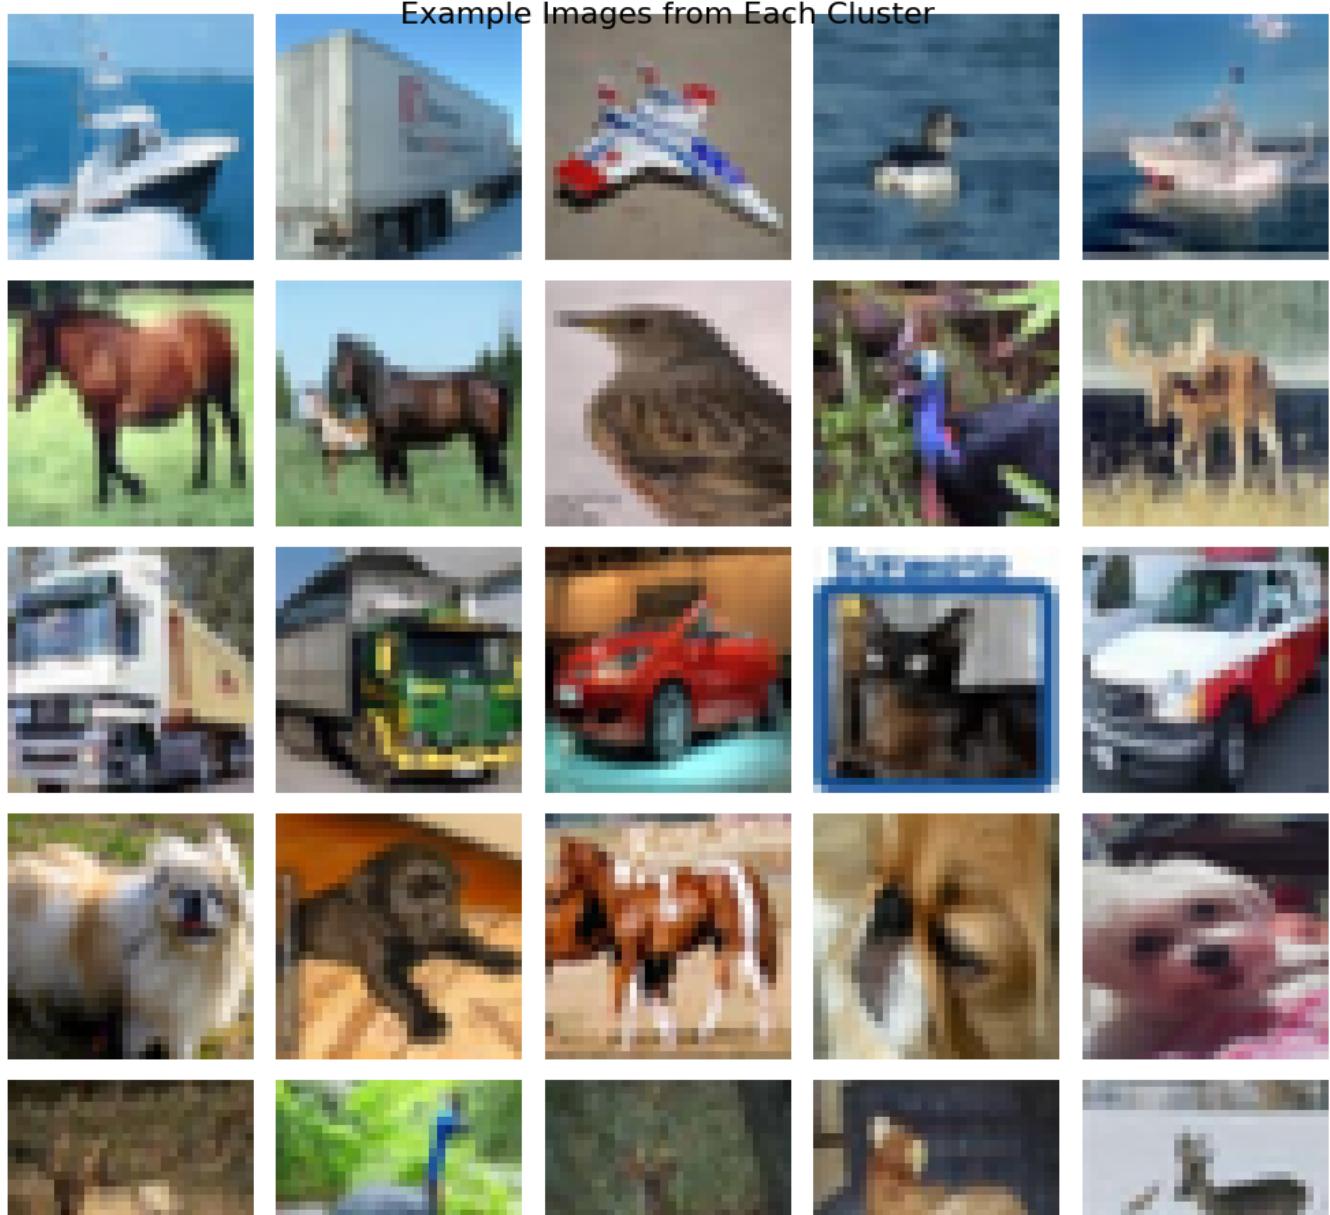
⌚ Adjusted Rand Index (ARI): 0.1408

t-SNE of VGG16 Features + KMeans Clustering
ARI: 0.1408





Example Images from Each Cluster



Start coding or generate with AI.

